



One Write.



One Hundred Million Reads.

AN ENGINEERING STORY ABOUT
HOW DISTRIBUTED SYSTEMS ACTUALLY EVOLVE

KUSH MISHRA

One Write. One Hundred Million Reads.

One Write. One Hundred Million Reads.

Copyright © Kush Mishra. All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

This is a work of engineering non-fiction. The systems and scenarios described are generalized composites intended to illustrate architectural patterns, and do not depict the specific implementation of any named company or product.

First edition.

How to Read This Book

Each chapter in this book begins with an incident, not a definition.

You will meet a system in trouble before you meet the concept that explains it. That is deliberate. The lesson tends to land harder once you have felt the failure first.

If you would rather know where a chapter is headed before you start it, the Table of Contents lists a short description under each title. And if you want to see the shape of the whole book before reading a single page, turn to “One Sentence from Every Chapter” near the end — it distills every chapter's lesson into one line, in order.

Either approach works. The chapters do not require the summary to make sense, and the summary does not require the chapters. Use whichever helps you read.

Contents

FRONT MATTER

Preface	11
Prologue	17

CHAPTERS

1.	The Simplest Backend You'll Ever Build	25
	<i>The baseline: one app, one database — no premature infrastructure</i>	
2.	Build the Wrong System First	35
	<i>Why over-engineering fails before under-engineering does</i>	
3.	Nothing Is Broken. So Why Is Everyone Complaining?	47
	<i>Connection pool exhaustion</i>	
4.	PostgreSQL Didn't Fail. You Did.	57
	<i>Missing indexes and the N+1 query problem</i>	
5.	The Most Expensive Question in Engineering	69

Deciding whether — and when — to introduce a cache

6.	The Fastest Database	79
	Query Costs Zero	
	<i>Introducing the cache-aside pattern</i>	
7.	Every Cache Eventually	89
	Lies	
	<i>Cache invalidation and staleness</i>	
8.	Your Biggest DDoS Attack	99
	Doesn't Know It's Attacking	
	You	
	<i>Thundering herd and request coalescing</i>	
9.	When One Server Becomes	109
	Two	
	<i>Horizontal scaling and shared state across servers</i>	
10.	The Decision Maker	121
	Nobody Notices	
	<i>Load balancing algorithms</i>	
11.	Yesterday's Truth	131
	<i>Stale health checks — the present is distributed</i>	
12.	The System Starts	141
	Watching Itself	
	<i>Introducing observability and metrics</i>	
13.	Four Witnesses, One	153
	Incident	
	<i>Correlating logs across distributed services</i>	

14.	The Waiting Game	163
	<i>Timeouts</i>	
15.	Hope Is Not a Strategy	173
	<i>Retries and backoff — and how naive retries worsen outages</i>	
16.	Learning to Say No	183
	<i>Circuit breakers</i>	
17.	When Everyone Decides	193
	<i>Leader election</i>	
18.	The Day Truth Split in Two	203
	<i>Split-brain</i>	
19.	When Nobody Gets to Decide Alone	213
	<i>Consensus and quorum</i>	
20.	The Memory Everyone Must Share	223
	<i>Introducing replication and shared state</i>	
21.	The Half-Written Story	233
	<i>Partial writes and atomicity</i>	
22.	The Diary	243
	<i>Write-ahead logging (WAL)</i>	
23.	The Cost of Forgetting	253
	<i>Durability and replication lag</i>	
24.	One Write	263
	<i>The single-writer principle and write ordering</i>	

25. One Hundred Million Reads 273

Scaling reads — the payoff, and the title, revealed

BACK MATTER

Tomorrow Morning, 285

Revisited

Epilogue 297

The Journey at a Glance 307

One Sentence from Every 311

Chapter

Preface

There is a certain kind of engineering problem that rarely gets discussed.

Not because it is uncommon.

But because it looks deceptively easy.

Imagine a backend that computes an answer exactly once.

Hours later, or perhaps days later, millions of users arrive asking only one question.

"What's my result?"

No writes.

No transactions.

No race conditions.

No distributed consensus.

Just reads.

On paper, this sounds like one of the simplest systems an engineer could build.

In production, it can become one of the hardest.

Over the years, I realized that many seemingly unrelated systems share this exact characteristic.

An IPO allotment portal.

A university entrance examination.

Election results.

A lottery.

Government application portals.

Sports tournament brackets.

Even some internal enterprise systems.

Different products.

Different users.

Different domains.

Yet they all face the same engineering challenge.

A single dataset.

An enormous, synchronized wave of readers.

Despite this pattern appearing repeatedly across industries, I found that most engineering resources discuss only the individual technologies involved.

There are books about databases.

Books about caching.

Books about networking.

Books about load balancing.

Books about distributed systems.

All excellent.

But very few walk through the engineering journey of *why* these systems evolve the way they do.

Why does a monolith stop being enough?

Why isn't the database the first thing to fail?

Why do retries make outages worse?

Why does adding a cache solve one problem while quietly creating another?

Why do systems accumulate complexity over time?

Those questions fascinated me far more than the technologies themselves.

This book is my attempt to answer them.

Not by presenting a finished architecture.

But by building one together.

We'll start with almost nothing.

One application.

One database.

One deployment.

No Redis.

No Kubernetes.

No message queues.

No globally distributed infrastructure.

Then we'll let reality intervene.

Every chapter introduces exactly one new failure.

Every architectural decision exists because the previous system could no longer survive.

Nothing is added because it's fashionable.

Nothing is added because every "modern architecture" includes it.

Every layer must earn its place.

The systems discussed throughout this book are intentionally generalized.

The examples are inspired by a class of real-world engineering problems that many organizations encounter, but the goal is not to document any specific implementation or company.

Instead, the focus is on developing an engineering mindset.

If, by the end of this book, you find yourself asking "*What problem is this technology actually solving?*" before introducing another component into your architecture, then this book has succeeded.

Engineering is rarely about choosing the most sophisticated solution.

More often, it is about resisting sophistication until reality demands it.

That's the journey we're about to begin.

Let's build the simplest system that could possibly work.

Then let's watch it fail.

Prologue

The Day the Internet Showed Up

It is 8:47 PM.

You are about to close your laptop after another ordinary day at work.

Then your phone vibrates.

Tomorrow at 9:00 AM, nearly one hundred million people may visit our platform within a few hours. We need to be ready.

You read the message twice.

One hundred million.

Not over a month.

Not over a year.

Not even over a week.

Tomorrow.

You open your monitoring dashboard.

Everything looks healthy.

CPU usage is below 20%.

Memory usage is stable.

Database latency is comfortably under ten milliseconds.

Error rates are practically zero.

If someone walked past your desk, they would probably congratulate you on running such a healthy backend.

But you know better.

Healthy systems fail every day.

Not because they were slow.

Not because they were buggy.

But because they were asked to do something they had never been asked to do before.

The internet has a peculiar habit.

Most days, it behaves politely.

Users arrive gradually.

Traffic follows predictable patterns.

Your autoscalers have time to react.

Your caches warm naturally.

Your databases settle into comfortable rhythms.

Then, occasionally, the internet decides to move as one.

An IPO allotment is published.

A national examination announces its results.

Election outcomes become official.

A government portal opens applications.

A lottery declares its winners.

Millions of people reach for the same answer at almost exactly the same moment.

They all ask different questions.

"Did I get allotted?"

"What's my rank?"

"Was my application accepted?"

"Did I win?"

Yet from an engineering perspective, they are all asking the same thing.

"Read my result."

Nothing is being updated anymore.

The difficult computation already happened hours—or even days—earlier.

Every answer already exists.

The system isn't busy calculating.

It is simply trying to survive being asked the same question millions of times.

At first glance, this seems like the easiest backend you could ever build.

No concurrent writes.

No distributed transactions.

No inventory races.

No payment consistency.

No conflict resolution.

Just read the data and send it back.

Simple.

Or at least, that's what almost everyone thinks.

History suggests otherwise.

Some of the largest production incidents on the internet have occurred not because writing data was difficult, but because reading it became overwhelming.

A backend that comfortably serves ten thousand users can collapse under ten million readers—not because the business logic changed, but because every assumption hidden inside the architecture quietly stopped being true.

Connections become scarce.

Caches become liabilities.

Retries multiply traffic.

Load balancers make imperfect decisions.

Networks become bottlenecks.

Systems that looked beautifully designed on Monday suddenly appear fragile on Tuesday.

The irony is almost cruel.

The data never changed.

Only the number of people who wanted to read it.

This book is about that moment.

Not about distributed systems in their entirety.

Not about every scaling problem ever encountered.

This is a story about a very particular class of systems.

Systems where the hardest work happens before anyone arrives.

Systems that write once.

Then spend the rest of their lives answering questions.

We will not begin with Kubernetes.

Or Redis.

Or Kafka.

Or CDNs.

Those technologies will appear eventually.

But only after they become necessary.

Because production systems do not evolve by following
architecture diagrams.

They evolve by surviving failure.

So we'll do the same.

We'll start with the smallest system that could possibly work.

Then we'll let reality break it.

Again.

And again.

Each failure will force exactly one architectural decision.

Nothing will be introduced early.

Nothing will be introduced simply because it is fashionable.

Every piece of complexity must earn its place.

By the end of this journey, you won't just know how to build a system capable of serving one hundred million readers.

You'll understand why every layer exists.

And more importantly—

you'll know when it doesn't.

Welcome.

Tomorrow, the internet arrives.

CHAPTER 1

The Simplest Backend You'll Ever Build

There is a dangerous sentence that engineers love to say.

"It's just a read API."

On the surface, they're right.

The system we're about to build has remarkably few responsibilities.

A user's result has already been computed.

It has already been validated.

It has already been stored.

Nothing changes after publication.

The backend performs exactly one operation.

A user asks for their result.

The server returns it.

That's it.

No shopping carts.

No payments.

No inventory reservations.

No collaborative editing.


No distributed transactions.

No background workflows.

No conflicting updates.

No race conditions.

If you were interviewing for a backend role and someone described this problem in two sentences, you would probably sketch a solution in less than a minute.



Congratulations.

You have just designed the architecture that this book will begin with.

And surprisingly—

it is the correct answer.

Not the beginner's answer.

Not the "temporary" answer.

The correct answer.

At least for now.

. . .

One of the most expensive habits in software engineering is solving tomorrow's problems with today's code.

Engineers love to imagine scale.

We picture millions of users before we've served a hundred.

We debate distributed databases before creating our first table.

We discuss global replication before deploying our first service.

The imagination of an engineer grows much faster than the traffic of a product.

Reality is usually far less dramatic.

Most systems never become internet-scale.

Many never outgrow a single machine.

Some of the most successful companies in the world spent years running architectures that would be dismissed in conference talks as "too simple."

Complexity is seductive because it feels like preparation.

Often, it is simply procrastination disguised as architecture.

Every component you introduce today becomes tomorrow's maintenance burden.

That burden doesn't disappear because the component was technically impressive.

Someone has to deploy it.

Monitor it.

Upgrade it.

Debug it.

Explain it to new engineers.

Wake up when it fails.

Every box on an architecture diagram quietly creates work for someone.

Usually the future version of you.

. . .

Let's define the problem more carefully.

Imagine an examination authority has finished evaluating answer sheets.

For every candidate, exactly one result now exists.

That result will never change.

Tomorrow morning, candidates will search using their application number.

The backend receives a request.

It retrieves exactly one row.

It returns exactly one response.

If we ignore scale for a moment, the engineering challenge appears almost laughably small.

The request path contains almost no business logic.

Receive a request.

Validate the input.

Read one record.

Return one response.

If this system served one hundred users tomorrow, how would it perform?

Probably perfectly.

If it served one thousand?

Probably perfectly.

Ten thousand?

Still likely perfectly.

So why would we build anything more complicated?

We shouldn't.

Not yet.

. . .

There is a temptation—especially among experienced engineers—to mistake sophistication for quality.

We've all seen architecture diagrams covered in colorful rectangles.

Microservices.

Caches.

Queues.

Service meshes.

Event buses.

API gateways.

Distributed tracing.

Circuit breakers.

Autoscaling groups.

Multi-region deployments.

Each component solves a real problem.

But there is an equally important question.

Does that problem exist today?

If the answer is no, then the component is not solving a problem.

It is creating one.

This is one of the quiet paradoxes of engineering.

The easiest system to operate is the one you never built.

. . .

Let's make one assumption.

Tomorrow, exactly one hundred people will use the system.

Not one hundred million.

One hundred.

How many application servers do we need?

One.

How many databases?

One.

How many caches?

None.

How many load balancers?

None.

How many Kubernetes clusters?

None.

How many regions?

One.

How many replicas?

Probably none.

At this scale, introducing any of those technologies would make the system objectively worse.

More code.

More moving parts.

More operational complexity.

More places to fail.

Without improving the user experience by even a millisecond.

This isn't minimalism for philosophical reasons.

It's engineering economics.

Every decision has a cost.

A good engineer doesn't ask,

"What can I add?"

A good engineer asks,

"What can I avoid adding?"

. . .

If someone asked you to draw this architecture on a whiteboard,
you could finish before they sat down.

And that's exactly how it should feel.

Simple.

Obvious.

Almost boring.

Don't underestimate that feeling.

Some of the best architectures in history looked boring when they
were young.

They only became interesting after reality forced them to evolve.

Our system hasn't reached that point yet.

It has no bottlenecks.

No outages.

No operational pain.

Nothing has earned the right to become more complicated.

So we stop here.

One application.

One database.

One purpose.

Tomorrow morning, one hundred people will check their results.

The system will work flawlessly.

Which raises an uncomfortable question.

If the architecture is correct...

why are we talking about scaling at all?

CHAPTER 2


Build the Wrong System First

Yesterday, our backend served one hundred users.

Nothing failed.

Every request completed in a few milliseconds.

The architecture was almost embarrassingly simple.



No caches.

No replicas.

No queues.

No orchestration.

No distributed anything.

It worked exactly as intended.

So today, your manager walks over to your desk.

"We're expecting ten thousand users tomorrow."

You smile.

Ten thousand isn't one hundred million.

Surely the architecture needs to evolve.

Right?

Not yet.

. . .

Software engineering has an unusual relationship with the future.

We spend an extraordinary amount of time preparing for events that may never happen.

Ask a room full of engineers to design a backend for an online bookstore.

Within minutes someone will suggest microservices.

Someone else will recommend Kafka.

Another engineer will propose CQRS.

Before a single customer has signed up, the architecture diagram resembles the infrastructure of a Fortune 500 company.

The irony is that almost every successful system started much smaller.

Google was once a handful of machines.

Facebook began in a college dormitory.

GitHub ran on a Ruby on Rails monolith for years.

Stack Overflow served millions of developers from a famously simple architecture long before anyone considered breaking it apart.

Complexity did not create their success.

Success forced complexity upon them.

That's a subtle but important distinction.

Scale is rarely designed. More often, it is inherited.

. . .

Let's return to our examination portal.

The results are already published.

Every request still follows the same path.

Receive an application number.

Look up one record.

Return one response.

The only thing that changed is the number of people asking.

Business logic is identical.

Database schema is identical.

API contract is identical.

The code didn't suddenly become slower overnight.

The workload simply became larger.

This distinction matters.

Many engineers instinctively respond to higher traffic by changing architecture.

Experienced engineers ask a different question.

"Which assumption actually stopped being true?"

Until you answer that question, every architectural change is speculation.

. . .

Suppose one request takes 8 milliseconds to complete.

That sounds excellent.

Now suppose ten thousand users arrive over the course of an hour.

Can one server handle that?

Probably.

What if they arrive over ten minutes?

Still possible.

Five minutes?

Maybe.

One minute?

Now we're asking a different question.

Notice what happened.

The problem wasn't the number of users.

It was **how quickly they arrived**.

Traffic is not merely a count.

Traffic has shape.

A million users spread across an entire day look very different from a million users arriving within sixty seconds.

Engineers often measure volume.

Production systems fail because of concentration.

A river is harmless until all of its water tries to pass through the same narrow bridge.

. . .

At this point, many teams make their first architectural mistake.

Someone says,

"We should split the service."

Into what?

The authentication endpoint?

The result endpoint?

The health check?

Nothing in our workload suggests those responsibilities need to live independently.

Microservices solve organizational and scaling problems.

We haven't encountered either.

Splitting a healthy application into five services doesn't remove complexity.

It redistributes it.

Now every request crosses the network.

Failures become distributed.

Deployments multiply.

Observability becomes harder.

Local debugging becomes slower.

You've solved a problem you never had.

And created several you definitely do.

. . .

Another engineer proposes Kubernetes.

Why?

"We'll need it eventually."

Perhaps.

But "eventually" is not a deployment strategy.

Every operational platform carries a learning curve.

Every abstraction introduces failure modes.

Clusters fail.

Schedulers fail.

Ingress controllers fail.

Certificates expire.

None of those risks existed yesterday.

They exist now because we added them.

Technology is never free.

Even when it's open source.

. . .

Someone else recommends introducing Redis.

Again, why?

Our database is serving requests comfortably.

Every cache introduces a new question.

How long should data remain cached?

What happens when the cache misses?

Who invalidates stale entries?

What happens if Redis becomes unavailable?

Caching is one of the greatest performance optimizations ever invented.

It is also one of the fastest ways to create consistency problems.

Until the database becomes a bottleneck, Redis isn't solving a problem.

It is patiently waiting to become one.

. . .

The uncomfortable truth is that engineering rewards restraint far more often than ambition.

There is no award for introducing the largest number of technologies.

Users don't care how many services you deployed.

They care whether the page loads.

Simple systems are easier to understand.

Easier to monitor.

Easier to recover.

Easier to explain.

Easier to replace.

Every additional component should justify its existence with measurable value.

Otherwise it is merely decoration.

Architecture is not the art of adding boxes. It's the discipline of refusing them.

. . .

So we resist the temptation.

The architecture remains exactly the same.

One application.

One database.

Nothing more.

Not because we're ignoring the future.

Because the future hasn't arrived yet.

And when it does, it will tell us precisely what to build next.

The question is whether we'll listen.

. . .

Tomorrow morning, the system goes live.

Ten thousand candidates begin checking their results.

The application server barely notices.

PostgreSQL remains calm.

CPU usage stays comfortably low.

Memory usage is stable.

Response times remain excellent.

By every metric available to us, the system is healthy.

And yet...

support tickets begin appearing.

Users complain that the website feels slow.

Some say it won't load at all.

Others report that refreshing eventually works.

The database isn't overloaded.

The backend isn't overloaded.

So why does the website still feel broken?

That answer isn't hiding inside our application.

It's waiting somewhere much closer to the user.

CHAPTER 3

Nothing Is Broken. So Why Is Everyone Complaining?

At 9:00 AM, the results go live.

For a brief, glorious moment, everything looks exactly as you hoped.

The deployment is healthy.

The application starts without errors.

Database connections are established.

Health checks are green.

CPU usage hovers around 18%.

Memory sits comfortably below 40%.

No alerts.

No exceptions.

No panic.

You lean back in your chair.

Maybe this really was just a read API.

• • •

At 9:07 AM, your first support ticket arrives.

"The website is taking forever to load."

Strange.

You refresh your monitoring dashboard.

Nothing unusual.

Response time: 11 ms.

Database latency: 6 ms.

Error rate: 0%.

Every graph says the same thing.

The system is healthy.

• • •

Another ticket appears.

"It keeps spinning."

Then another.

"Sometimes it works after refreshing."

Five minutes later, your team's Slack channel becomes unusually active.

Product managers begin forwarding screenshots from social media.

People aren't complaining about incorrect results.

They're complaining that they can't see them.

Yet every dashboard insists everything is fine.

Which one do you trust?

The users...

or the graphs?

• • •

This is one of the first uncomfortable lessons of production engineering.

A healthy backend does not guarantee a healthy user experience.

Those are two very different statements.

And confusing them has wasted countless hours inside incident calls.

. . .

You open your browser.

Enter an application number.

Press Enter.

The page loads.

Eventually.

It feels sluggish.

Not catastrophically slow.

Just...

hesitant.

You open the Network tab.

For the first time today, you're no longer looking at your backend.

You're looking at your user.

Immediately, something feels different.

The first request doesn't even reach your application.

The browser is still busy.

DNS lookup.

TCP connection.

TLS negotiation.

Downloading HTML.

Downloading CSS.

Downloading JavaScript.

Loading fonts.

Fetching icons.

Only then...

does your API request finally leave the browser.

Your backend spent eleven milliseconds answering the request.

The browser spent hundreds of milliseconds simply preparing to ask it.

Suddenly, those beautiful backend metrics don't feel quite so impressive.

. . .

Engineers naturally measure the systems they build.

Users measure the systems they experience.

Those are rarely the same thing.

Imagine walking into a restaurant.

The chef prepares your meal in three minutes.

Excellent.

Unfortunately, you waited twenty-five minutes before anyone took your order.

When you leave, do you praise the chef's efficiency?

Probably not.

You tell your friends,

"The restaurant was slow."

Technically, that isn't true.

Practically, it is.

Production systems are judged by end-to-end experience, not individual components.

. . .

You decide to investigate further.

You compare timestamps.

The backend receives the request at 09:11:23.412.

It finishes at 09:11:23.423.

Eleven milliseconds.

The browser renders the completed page at 09:11:24.087.

Nearly seven hundred milliseconds later.

Where did the rest of the time go?

Not inside your code.

Not inside PostgreSQL.

Not inside your business logic.

It disappeared into everything surrounding them.

Networks.

Protocols.

Browsers.

Static assets.

Encryption.

Distance.

The internet itself.

. . .

This realization changes the way you think about performance.

Performance isn't a property of a server.

It's a property of a journey.

Every request travels farther than most engineers imagine.

Before your application executes a single line of business logic,
dozens of invisible systems have already done work on its behalf.

Most users don't know they exist.

Most engineers forget they exist.

Until they become the bottleneck.

***Your API isn't the beginning of a request. It's
merely one stop along the way.***

. . .

By lunchtime, the incident is over.

Not because you changed the backend.

Because someone enabled long-lived caching for static assets, compressed JavaScript bundles, and fixed an inefficient frontend loading sequence.

The API never changed.

The database never changed.

Users simply reached them faster.

. . .

That evening, you write one sentence in your notebook.

"Today I learned that optimizing the fastest part of a request is rarely the fastest way to improve it."

You don't know it yet.

But tomorrow, you're going to encounter the opposite problem.

The browser will become fast.

The network will become fast.

The backend will become fast.

And for the first time...

the database will finally have nowhere left to hide.

CHAPTER 4

PostgreSQL Didn't Fail. You Did.

Three days later, the traffic pattern changes.

Not dramatically.

Just enough.

The morning spike grows larger.

More schools publish their results.

More users arrive together.

The system still feels stable.

Until 9:14 AM.

An alert appears.

HTTP 503
Service Unavailable

A few seconds later.

Another.

Then dozens.

The incident channel comes alive.

"Database is down?"

Someone checks.

"No."

"Application crashed?"

"No."

"CPU?"

"32%."

"Memory?"

"Fine."

Everything looks... normal.

Except users can't access their results.

. . .

You log into PostgreSQL.

It's alive.

No restart.

No crash.

No corruption.

In fact, it's answering queries almost instantly.

The query itself isn't slow.

```
SELECT *  
FROM results  
WHERE application_number = ?;
```

POSTGRESQL DIDN'T FAIL. YOU DID.

Execution time:

3 milliseconds.

That can't possibly explain a failed request.

So where did it go?

. . .

You open another dashboard.

One you've rarely looked at before.

Database Connections.

Current connections:

500 / 500

Every slot is occupied.

Every new request waits.

Eventually, it gives up.

The database isn't overloaded.

The waiting room is.

. . .

This surprises almost everyone the first time they see it.

We instinctively think of databases as engines that execute queries.

But before a query can execute...

it needs permission to enter.

A connection.

No connection.

No query.

No result.

. . .

Imagine a restaurant with fifty tables.

The kitchen can prepare a meal in two minutes.

Fast enough.

But one hundred customers arrive simultaneously.

The kitchen isn't the first problem.

The dining room is.

There are no tables left.

The chef is idle.

Customers are waiting outside.

When they leave frustrated, nobody says,

"The tables were the bottleneck."

They simply say,

"The restaurant was full."

Databases behave the same way.

. . .

You inspect the application.

Every incoming request opens a new database connection.

Runs one query.

Closes the connection.

It seemed reasonable during development.

With one hundred users...

nobody noticed.

With thousands...

opening and closing connections becomes expensive.

Authentication.

Network negotiation.

Resource allocation.

Repeat.

Thousands of times.

Per minute.

You've optimized the query.

But you've ignored everything required before the query can even begin.

. . .

The first fix isn't a faster database.

It's fewer conversations with it.

You introduce a connection pool.

Instead of creating a new connection for every request, the application keeps a small group of already-established connections ready to work.

The requests line up.

The connections don't.

Suddenly, the graph changes.

Connection creation drops almost to zero.

Response times stabilize.

Error rates disappear.

The database didn't become faster.

It simply stopped repeating unnecessary work.

The fastest connection is the one you never had to create.

. . .

The incident ends.

But one graph still bothers you.

Average query time has slowly increased.

Not seconds.

Milliseconds.

Enough to notice.

You examine the execution plan.

The query is scanning the entire table.

Every request.

Every user.

Every time.

Why?

Because the database has no efficient way to find an application number.

You've asked it to search an entire phone book every time someone asks for one name.

The solution isn't magical.

It's an index.

Tomorrow morning, PostgreSQL no longer searches.

It navigates.

The difference is almost invisible in development.

In production, it's transformational.

Not because indexes are clever.

Because unnecessary work compounds.

One extra millisecond feels irrelevant.

Multiply it by ten million requests.

Now you've spent nearly three hours of database time doing work that never needed to exist.

. . .

This is the first time our architecture genuinely improved.

Not because we replaced PostgreSQL.

Not because we introduced another service.

Because we understood where time was actually being spent.

Good engineering rarely begins with optimization.

It begins with observation.

. . .

That afternoon, someone asks a familiar question.

"Should we finally add Redis?"

You pause.

For the first time...

the question doesn't sound unreasonable.

The database is healthy.

Connections are under control.

Queries are indexed.

Execution plans look excellent.

And yet...

every request still travels all the way to PostgreSQL.

To retrieve data that hasn't changed once all day.

Millions of identical reads.

The same journey.

Again.

Again.

Again.

If the answer never changes...

why are we asking the database every single time?

CHAPTER 5

The Most Expensive Question in Engineering

The meeting begins at 2:00 PM.

The morning incident is over.

The graphs have settled.

Users are happy again.

Someone shares the incident summary.

- Added a connection pool.
- Created the missing index.
- Reduced database wait times.
- Error rate back to zero.

Everyone nods.

The system survived.

Then someone asks a question that instantly changes the room.

"Should we add Redis?"

Silence.

Not because nobody knows Redis.

Quite the opposite.

Everyone does.

Which makes the question dangerous.

. . .

Good engineering meetings rarely begin with disagreement.

They begin with enthusiasm.

Redis is fast.

Ridiculously fast.

It lives in memory.

It can answer requests in microseconds.

We've all seen the conference talks.

We've all seen the benchmark graphs.

We've all seen the architecture diagrams.

Application.

Redis.

Database.

It almost feels incomplete without that middle box.

So why are we hesitating?

. . .

Because technology is very good at solving the problems we imagine.

Production is much better at revealing the problems we actually have.

Those are not always the same.

. . .

You walk to the whiteboard.

Instead of drawing Redis...

you draw a single request.



Then you ask the room.

"What part of this journey is actually hurting us?"

Nobody answers immediately.

Someone finally says,

"The database."

You shake your head.

"No."

Another engineer replies,

"The query."

"No."

Someone else suggests,

"Disk I/O?"

"Still no."

The room grows quieter.

You point to the morning's graphs.

The database wasn't slow.

The query wasn't slow.

The CPU wasn't busy.

The memory wasn't exhausted.

Connections were.

And we already solved that.

So why are we introducing a completely new system?

. . .

This is one of the easiest traps in engineering.

Mistaking familiarity for necessity.

Redis is familiar.

So are Kafka.

Kubernetes.

ElasticSearch.

NGINX.

Prometheus.

They're excellent tools.

None of them become excellent merely because they exist.

They become excellent when a measurable problem justifies their cost.

***Every technology enters your architecture
carrying an invoice.***

Sometimes the invoice is money.

More often...

it's complexity.

Someone must monitor Redis.

Back it up.

Upgrade it.

Patch it.

Recover it.

Document it.

Explain it to every new engineer.

Wake up when it goes down.

The benchmark graph never shows those costs.

. . .

So instead of adding Redis...

you ask another question.

"What happens if PostgreSQL disappears?"

The room laughs.

"It can't."

"Imagine it does."

Now the conversation changes.

How would users receive results?

They wouldn't.

Interesting.

Redis wouldn't save us either.

Because Redis isn't a source of truth.

It's a copy.

Suddenly everyone realizes something important.

Caching isn't about replacing the database.

It's about avoiding unnecessary conversations with it.

That distinction matters.

. . .

You leave the meeting without adding Redis.

Some people feel disappointed.

It feels like inaction.

It isn't.

Restraint is often invisible.

Nobody celebrates the technologies you wisely refused to introduce.

They only notice the outages you avoided because you kept the system understandable.

. . .

A week passes.

Traffic grows again.

The database remains healthy.

Connections remain healthy.

Queries remain healthy.

Everything remains...

healthy.

Until one morning, you notice something curious.

The monitoring dashboard reports nearly twelve million database reads.

Not writes.

Reads.

Every single one returns information that was generated days ago.

The rows haven't changed.

Not once.

Twelve million identical journeys.

Twelve million identical answers.

Twelve million identical conversations with the same database.

You stare at the graph.

Then at the application.

Then back at the graph.

For the first time since this project began...

Redis isn't solving an imagined problem anymore.

It's solving a measurable one.

And that changes everything.

CHAPTER 6

The Fastest Database Query Costs Zero

By now, our architecture has survived several lessons.

We learned that a simple system is often the correct system.

We learned that users experience journeys, not APIs.

We learned that databases can become constrained long before they become slow.

Each improvement solved a real problem.

Each improvement earned its place.

Today, we're about to introduce the first entirely new component into our architecture.

Not because Redis is fashionable.

Because PostgreSQL has become... repetitive.

• • •

Open the production dashboard.

Ignore the CPU graphs.

Ignore the memory graphs.

Ignore the error rates.

Instead, look at a single number.

Database Reads Today

12,184,391

Now ask yourself a different question.

How many of those reads were actually necessary?

Not successful.

Necessary.

The answer is uncomfortable.

Almost none.

. . .

Every request asks exactly the same thing.

"What is my result?"

The application receives the request.

Finds the candidate.

Looks up one row.

Returns one response.

Five seconds later...

another request arrives.

Same candidate.

Same row.

Same answer.

Five minutes later...

another.

The answer hasn't changed.

It can't change.

The examination board published the results three days ago.

The row is identical.

Yet the application insists on walking all the way to PostgreSQL every single time.

Imagine visiting your local library.

Every morning, you borrow the exact same book.

You read the same page.

You return it.

Tomorrow, you repeat the journey.

Eventually someone asks,

"Why don't you just keep the book on your desk?"

Caching begins with that question.

Not with Redis.

. . .

There is a subtle misunderstanding that follows almost every discussion about caches.

People say,

"A cache makes things faster."

That's true.

But incomplete.

A cache does something much more interesting.

It changes **where** work happens.

Instead of asking the database repeatedly,

the application remembers.

It answers from memory.

The database remains the source of truth.

The application simply stops bothering it with questions whose answers it already knows.

The fastest query isn't the optimized one. It's the one you never execute.

. . .

Notice what we are *not* doing.

We're not moving authority.

If a result changes, PostgreSQL is still correct.

If Redis disappears, PostgreSQL is still correct.

If every cache entry vanishes, PostgreSQL is still correct.

A cache is confidence.

Not truth.

That distinction will save you from countless production incidents.

. . .

You decide to perform a small experiment.

Instead of asking PostgreSQL every time,

the application checks memory first.

If the answer exists,

return it immediately.

If it doesn't,

ask PostgreSQL,

return the result,

and remember it for next time.

The next morning, traffic arrives.

The graph looks different.

Very different.

Database reads begin climbing.

Then...

they almost stop.

The application is still answering requests.

Millions of them.

But PostgreSQL receives only a tiny fraction.

The database didn't become faster.

It became quieter.

. . .

The operations dashboard now tells two stories.

One graph measures requests served.

It rises almost vertically.

Another graph measures database queries.

It barely moves.

For the first time since this journey began,

those two graphs are no longer the same.

The application has learned something.

Memory is often cheaper than repetition.

. . .

The incident review is almost disappointingly short.

No connection exhaustion.

No query tuning.

No emergency indexing.

Just one observation.

"We stopped asking the same question millions of times."

Sometimes engineering breakthroughs aren't about building faster systems.

They're about building less forgetful ones.

. . .

That afternoon, a new engineer joins the team.

They study the architecture diagram.

Application.

Redis.

PostgreSQL.

They ask an innocent question.

"If PostgreSQL already has the data...

why keep another copy?"

The room smiles.

Because every engineer asks that question once.

And every engineer eventually learns the same answer.

Copies are easy.

Keeping copies correct...

is where the real engineering begins.

Tomorrow, our data won't change.

But someday it will.

And on that day,

our beautifully quiet cache will become the source of an entirely new class of production incidents.

CHAPTER 7

Every Cache Eventually Lies

For nearly three weeks, nobody talks about Redis.

Which, strangely enough, is the greatest compliment a cache can receive.

The incident channel is quiet.

Database graphs have flattened.

Morning traffic arrives.

Morning traffic leaves.

Users receive their results.

Engineers sleep through the night.

Life is good.

Then a support ticket arrives.

It doesn't sound urgent.

"My result changed."

You read it once.

Then again.

Impossible.

Results are immutable.

That's the entire premise of this system.

The result cannot change.

You check PostgreSQL.

The row hasn't changed.

You check the audit logs.

Nothing.

No updates.

No corrections.

No deployments.

No data migration.

Everything says the same thing.

The database is telling the truth.

So why is the user seeing something else?

• • •

You open the application.

Search for the same candidate.

The result appears immediately.

Exactly as the user described.

Wrong.

Refresh.

Still wrong.

A minute later...

it's correct.

No code changed.

No database update.

Nothing happened.

Except time.

• • •

This is the first time you've caught your cache in a lie.

Not a malicious lie.

Not even a software bug.

Just...

an old memory.

. . .

Human memory works the same way.

Imagine someone asks for your home address.

You answer instantly.

You don't search your email.

You don't open Google Maps.

You remember.

Most of the time...

that's exactly what you want.

Unless you've moved.

Now your memory is fast.

Confident.

And completely wrong.

Caches make the same trade.
They exchange certainty for speed.
Usually that's an excellent bargain.
Until reality changes underneath them.

Every cache answers two questions.

"Do I know the answer?"

"Can I still trust it?"

The first question is easy.
The second is the entire discipline.

. . .

The team gathers around the whiteboard.
Someone suggests reducing the cache lifetime.
Another suggests clearing Redis every hour.
Someone proposes restarting the cache every morning before
results are published.
Each idea sounds reasonable.
Each one quietly admits the same thing.

Nobody knows exactly when the data stops being trustworthy.

. . .

Then someone asks a question nobody expected.

"When was the last time our data actually changed?"

The room becomes silent.

Three weeks ago.

Before Redis was even introduced.

The cache isn't wrong because it held stale information.

The cache is wrong because one engineer manually corrected a candidate's result yesterday afternoon.

One row.

One exceptional case.

One emergency fix.

PostgreSQL was updated.

Redis wasn't.

The source of truth changed.

The copy never heard about it.

. . .

This is the hidden tax of every cache.

Reading becomes cheaper.

Consistency becomes more expensive.

The database only worries about storing data.

A cache worries about remembering when to forget.

. . .

You fix the immediate problem.

The application removes the outdated entry.

The next request rebuilds it.

Users receive the correct answer.

The incident ends.

But the architecture has changed.

Not physically.

Mentally.

Nobody in the room refers to Redis as "the cache" anymore.

They start calling it

"the copy."

It's a tiny change in language.

A profound change in thinking.

Copies drift.

Copies expire.


Copies disagree.

Copies need maintenance.

Truth doesn't.

. . .

That evening, you stare at the architecture diagram.



Yesterday it looked simple.

Today it looks different.

Not because another box appeared.

Because one of those boxes now contains uncertainty.

For the first time, your system has two places capable of answering the same question.

Sometimes they agree.

Sometimes they don't.

Engineering just became less about performance...
and more about trust.

. . .

Weeks pass.

Traffic continues to grow.

Redis absorbs millions of requests.

PostgreSQL remains almost untouched.

The architecture feels mature.

Balanced.

Predictable.

Then something happens that nobody anticipated.

Not a database failure.

Not a cache failure.

A human failure.

One frustrated candidate presses Refresh.

Again.

Again.

Again.

Within seconds...

ten thousand people are doing exactly the same thing.

The database barely notices.

Redis barely notices.

The network certainly notices.

And suddenly you realize something unsettling.

The greatest threat to your system isn't stale data.

It's perfectly impatient users.

CHAPTER 8

Your Biggest DDoS Attack Doesn't Know It's Attacking You

The incident begins with a single refresh.

One candidate enters an application number.

The page takes slightly longer than usual.

Nothing dramatic.

Perhaps an extra second.

They do what every internet user has learned to do.

They press **Refresh**.

The page still doesn't load.

They press it again.

Then again.

Then they open another browser.

Then they try from their phone.

By the time the page finally appears, the candidate has unknowingly generated eight requests.

Not because they were impatient.

Because they were human.

. . .

At 9:18 AM, your dashboard reports something strange.

Traffic has doubled.

Marketing didn't launch a campaign.

No new examination results were published.

No notification emails were sent.

The number of users hasn't changed.

So why did the request rate suddenly double?

. . .

You compare two graphs.

The first shows unique visitors.

It remains almost perfectly flat.

The second shows HTTP requests.

It rises sharply.

The curves no longer resemble one another.

Somehow...

your users multiplied without multiplying.

. . .

Production has a peculiar habit of exposing assumptions you never realized you made.

One of those assumptions is surprisingly common.

One user equals one request.

Almost never.

One user might generate ten requests.

Or fifty.

Or hundreds.

A browser loads HTML.

Then CSS.

Then JavaScript.

Then fonts.

Then icons.

Then API calls.

Then analytics.

Then health checks.

Then retries.

Then the user presses Refresh.

The application isn't serving people.

It's serving behavior.

• • •

You replay a session recording.

A candidate clicks the search button.

Nothing happens immediately.

Three seconds pass.

The browser displays no feedback.

The candidate clicks again.

The application receives another identical request.

Then another.

Then another.

Each request is perfectly valid.

Each request deserves an answer.

Collectively...

they are overwhelming the system.

. . .

This is what makes these incidents so unsettling.

Nobody is doing anything wrong.

Nobody is attacking you.

Nobody is trying to cause harm.

Your users are simply responding to uncertainty.

Computers wait.

Humans don't.

Every loading spinner is a negotiation with human patience.

Lose that negotiation...

and users begin generating traffic you never planned for.

. . .

The backend behaves exactly as instructed.

Every request reaches Redis.

Redis returns the same result.

The application sends the same response.

Again.

Again.

Again.

Millions of times.

Technically...

every request is successful.

Operationally...

they're waste.

. . .

Someone suggests increasing server capacity.

Another proposes a larger Redis cluster.

A third recommends scaling PostgreSQL.

Those ideas might reduce today's symptoms.

They don't answer tomorrow's question.

Why are we doing the same work repeatedly for the same person?

. . .

You search the logs for one application number.

The result surprises everyone.

```
09:18:12 GET /result/487231
09:18:13 GET /result/487231
09:18:14 GET /result/487231
09:18:14 GET /result/487231
09:18:15 GET /result/487231
09:18:16 GET /result/487231
09:18:18 GET /result/487231
```

Seven identical requests.

Six identical answers.

Zero new information.

The system wasn't under heavy demand.

It was under repeated demand.

There is a difference.

. . .

That afternoon, you introduce something surprisingly small.

Not another database.

Not another cache.

Not another service.

Just an idea.

When a request for a candidate is already being processed...

don't start another one.

Let the others wait.

The first request does the work.

Everyone else receives the same answer.

For the user...

nothing changes.

For the infrastructure...

everything changes.

Seven requests become one.

Not because fewer people visited.

Because the system finally recognized that identical work shouldn't happen simultaneously.

***Scale isn't always about serving more requests.
Sometimes it's about refusing to repeat yourself.***

. . .

The graphs settle.

Unique visitors remain exactly the same.

Request volume falls dramatically.

CPU usage drops.

Redis traffic drops.

Network bandwidth drops.

The users never notice.

Which is precisely the point.

The best performance optimization is often invisible.

. . .

That evening, you write one sentence in your notebook.

"Today I learned that traffic isn't created by users.

It's created by uncertainty."

You close your laptop.

Tomorrow's challenge won't come from browsers.

Or databases.

Or caches.

It will come from something far more difficult to predict.

Geography.

The requests you've been serving all morning have one thing in common.

They've all been coming from the same city.

Tomorrow...

they won't.

CHAPTER 9

When One Server Becomes Two

Three months ago, this architecture fit on a whiteboard.

One application.

One database.

One engineer.

Now the whiteboard is full.

Redis.

Connection pools.

Indexes.

Request deduplication.

Monitoring dashboards.

Nothing was added because it looked impressive.

Everything was added because reality demanded it.

The architecture didn't grow.

It adapted.

. . .

Traffic continues to climb.

Not dramatically.

Steadily.

Every week looks almost identical to the last.

Five percent more users.

Eight percent more requests.

A few more support tickets.

A few more graphs.

Growth is surprisingly quiet.

Until it isn't.

. . .

One Tuesday morning, something feels different.

Not broken.

Busy.

CPU hovers around ninety percent.

Garbage collection becomes more frequent.

Response times slowly creep upward.

Nothing explodes.

The server simply looks...

tired.

. . .

Someone asks the obvious question.

"Should we optimize the code?"

Probably.

Someone else asks another.

"Should we buy a larger machine?"

Maybe.

A third engineer asks something much more interesting.

"What happens after the larger machine?"

The room becomes quiet.

Because everyone knows the answer.

Eventually...

there isn't a larger machine.

. . .

For months we've been solving problems by making one server better.

Better indexes.

Better queries.

Better caching.

Better connection management.

All excellent improvements.

But every one of them shares the same hidden assumption.

There is still only one application.

Today...

that assumption expires.

. . .

The easiest solution appears obvious.

Start another application.

The code is already stateless.

The API doesn't store user sessions.

Redis already holds temporary data.

PostgreSQL remains the source of truth.

Nothing prevents another copy from running.

Except one question.

If two servers now exist...

how does a browser decide which one to talk to?

. . .

You start both servers.

Then, almost absentmindedly, you refresh the page.

Server A.

Refresh.

Server B.

Refresh.

Server A again.

Without realizing it, you've just introduced something completely new into the system.

Choice.

Until yesterday, every request had exactly one destination.

Now...

every request has several.

Someone must decide.

Not occasionally.

Every single request.

Millions of times per day.

• • •

That decision sounds trivial.

Surely we can just alternate.

First request.

Server A.

Second request.

Server B.

Third request.

Server A again.

Simple.

Until one server becomes slower.

Or one begins restarting.

Or one receives a burst of long-running requests.

Suddenly, equal distribution no longer means equal work.

The architecture has acquired a new responsibility.

Not serving requests.

Directing them.

. . .

Notice something subtle.

Adding another application server didn't make the original server faster.

It didn't reduce latency.

It didn't improve database performance.

It simply created another place capable of doing the work.

Scale isn't multiplication.

It's coordination.

Every new server reduces individual work while increasing collective complexity.

That's the trade.

Always.

. . .

The first day ends successfully.

Both servers remain healthy.

Traffic is evenly distributed.

CPU falls below fifty percent.

Everyone celebrates.

Then someone asks a question that nobody had considered.

"What happens if Server A dies?"

The room answers confidently.

"Server B keeps serving traffic."

Correct.

Then another question.

"What happens to the requests already on Server A?"

Silence.

Some users retry.

Some browsers refresh.

Some requests disappear forever.

For the first time, availability isn't a property of a single machine.

It's a property of the system.


And systems fail differently than servers do.

. . .

That evening, your architecture diagram changes again.

Not because another box appeared.

Because another relationship did.



The applications still answer requests.

Redis still remembers.

PostgreSQL still tells the truth.

But there's a new actor in the story.

One that doesn't own data.

Doesn't execute business logic.

Doesn't answer users.

It simply decides.

Over.

And over.

And over again.

Tomorrow...

we're going to meet it.

And discover that making decisions for millions of requests is much harder than it looks.

ENGINEER'S NOTEBOOK

The first server taught us how to build software.

The second server teaches us how to build systems.

Those sound similar.

They aren't.

Software answers questions.

Systems decide where the questions should go.

CHAPTER 10

The Decision Maker Nobody Notices

Nobody deploys a load balancer because they want one.

They deploy it because the alternative has quietly stopped working.

Yesterday, every request had exactly one destination.

There was no decision to make.

A browser sent a request.

The application answered it.

Simple systems are full of certainties.

Distributed systems replace certainties with choices.

. . .

This morning, there are two application servers.

Both healthy.

Both running identical code.

Both connected to the same Redis instance.

Both reading from the same PostgreSQL database.

From the user's perspective, nothing has changed.

The URL is the same.

The page looks the same.

The response is the same.

Somewhere inside your infrastructure, however, a new question appears for every incoming request.

Which server should answer this one?

Not eventually.

Not every minute.

Every single request.

Millions of times.

. . .

At first, the answer seems obvious.

Alternate.

Server A.

Then Server B.

Then Server A again.

Engineers love symmetry.

Reality doesn't.

. . .

Around 10:30 AM, one of your application servers begins a full garbage collection cycle.

It pauses for a fraction of a second.

Not long enough to trigger an alert.

Long enough to delay a few requests.

The traffic distributor doesn't know that.

It continues sending every other request to the paused server.

Half your users experience slightly slower responses.

The other half don't.

The architecture did exactly what you told it to do.

The problem is...

you told it too little.

. . .

A week later, the team changes strategies.

Instead of alternating requests, send them to the server currently doing the least work.

That sounds smarter.

Usually it is.

Until one request unexpectedly takes five seconds.

Every new request avoids that server.

The second server becomes overloaded.

The first server eventually finishes.

Now it's almost idle.

The traffic swings back.

The system begins chasing its own shadow.

You solved one imbalance.

You created another.

. . .

This is one of the quiet truths of distributed systems.

There is rarely a perfect decision.

Only decisions that fail differently.

Every routing strategy is an opinion about the future.

Round robin assumes every request costs roughly the same.

Least-connections assumes current load predicts future load.

Latency-based routing assumes yesterday's fastest server will still be the fastest.

Sometimes those assumptions hold.

Sometimes they don't.

Engineering isn't choosing perfection.

It's choosing which assumption you're most willing to believe.

. . .

By lunchtime, something else becomes obvious.

The traffic distributor has become part of your architecture.

Not an accessory.

Not infrastructure someone else owns.

A genuine production dependency.

If it fails...

nobody reaches your applications.

Ironically, the component that performs almost no business logic now has the power to stop all business logic.

That realization changes how you think about it.

Yesterday it was just another box in a diagram.

Today it's part of your system.

. . .

You begin monitoring it.

Not because it became slower.

Because it became important.

How many requests is it forwarding?

How quickly?

How often does it retry?

What happens when an application stops responding?

How quickly does it notice?

Questions you never had to ask before now become operational necessities.

Growth has a habit of inventing entirely new responsibilities.

. . .

Late that afternoon, one application server crashes.

Nothing dramatic.

A bad deployment.

It happens.

The users barely notice.

New requests flow to the remaining server.

The incident lasts less than a minute.

The deployment is rolled back.

Everything returns to normal.

The team celebrates.

Then someone quietly asks,

"How did the traffic distributor know that server had failed?"

Nobody answers immediately.

Because nobody actually knows.

The server disappeared.

The requests stopped flowing.

Some component made a decision.

But what evidence did it use?

How long did it wait?

How many users received errors before that decision was made?

You realize something unsettling.

The hardest part of distributing traffic isn't choosing where to send requests.

It's deciding when to stop.

. . .

That evening, your architecture diagram changes again.

Not visually.

Conceptually.

You stop calling it a "load balancer."

You start calling it a **decision engine**.

The name is less familiar.

The idea is more accurate.

It doesn't balance load.

It continuously forms opinions about a changing system.

Sometimes those opinions are correct.

Sometimes they are outdated before they're even made.

Tomorrow...

we're going to discover what happens when those opinions depend on information that's already stale.

Because in distributed systems, information doesn't become dangerous when it's wrong.

It becomes dangerous when it's **old**.

ENGINEER'S NOTEBOOK

Adding another server didn't make the system distributed.

*Adding another **decision** did.*

Every time we introduce choice into an architecture, we also introduce uncertainty.

Computers execute instructions.

Distributed systems spend their lives making guesses.

CHAPTER 11

Yesterday's Truth

The deployment looked perfect.

Two application servers.

Traffic flowing evenly.

Redis answering millions of requests.

PostgreSQL almost idle.

The incident dashboard had been quiet for weeks.

Nobody expected Tuesday morning to be interesting.

It was.

. . .

At exactly 9:02 AM, one application server became unreachable.

Not permanently.

Just long enough.

Perhaps the network dropped a few packets.

Perhaps a switch restarted.

Perhaps the operating system paused for reasons nobody would ever completely understand.

For twenty-three seconds...

Server B disappeared.

Server A continued serving traffic.

No alarms.

No panic.

Twenty-three seconds later...

Server B returned.

Healthy.

Responsive.

Ready for work.

Except...

it wasn't receiving any.

. . .

You open the dashboard.

Server A is struggling.

CPU above ninety percent.

Response times rising.

Server B is almost idle.

Five percent CPU.

Barely doing anything.

Impossible.

The traffic distributor knows there are two healthy servers.

Doesn't it?

. . .

You inspect the health status.

Server A.

Healthy.

Server B.

Healthy.

Both green.

Then why is one drowning while the other waits?

. . .

Someone says,

"Restart the load balancer."

Before anyone moves...

another engineer asks,

"Why would that help?"

Nobody answers.

Because nobody really knows.

It has simply become one of those things engineers say during incidents.

. . .

Instead...

you start asking different questions.

When did the traffic distributor last check Server B?

How often does it ask?

What happens between those checks?

The answers are surprisingly ordinary.

Every few seconds...

the traffic distributor asks each application,

"Are you healthy?"

If the answer is yes...

traffic continues.

If the answer is no...

traffic stops.

Simple.

Elegant.

Reasonable.

Until reality becomes faster than your questions.

. . .

Imagine driving using yesterday's traffic report.

It isn't wrong.

It simply isn't current.

The highway was empty yesterday.

Today there's an accident.

Your decision was based on truth.

Just not today's truth.

Distributed systems suffer from the same problem.

Computers don't make decisions using reality.

They make decisions using the most recent information they happen to possess.

Those are very different things.

Every distributed decision is made with slightly outdated knowledge.

The moment you separate two computers...

that sentence becomes unavoidable.

. . .

The traffic distributor wasn't broken.

It simply believed Server B was still unavailable.

Because the last successful health check said so.

By the time the next health check arrived...

thousands of requests had already been sent elsewhere.

Nobody lied.

Nobody made a mistake.

Information simply aged.

. . .

This is one of the first principles that truly separates distributed systems from single-machine software.

Inside one process...

every function sees the same memory.

The same variables.

The same present.

Across machines...

there is no single present.

Only overlapping versions of it.

Some newer.

Some older.

Some already obsolete.

. . .

You begin looking around the architecture.

Suddenly...

you see the same pattern everywhere.

Redis contains yesterday's copy until it learns otherwise.

Application metrics arrive several seconds late.

Monitoring dashboards aggregate one-minute windows.

Autoscalers react to historical CPU usage.

Even your incident alerts describe problems that have already happened.

The entire system is quietly making decisions using history.

Not because it's poorly designed.

Because instantaneous knowledge doesn't exist.

. . .

That realization changes something fundamental.

You stop asking,

"Is this information correct?"

You start asking,

"How old is this information?"

Those questions sound similar.

They aren't.

One measures accuracy.

The other measures confidence.

. . .

By afternoon, the traffic distributor begins sending requests back to both servers.

The graphs return to normal.

The incident closes.

Nothing dramatic happened.

No outages.

No corruption.

No data loss.

Just twenty-three seconds of outdated information.

Twenty-three seconds were enough.

. . .

You leave the office thinking about something uncomfortable.

Every optimization we've introduced has relied on information.

Redis remembers previous reads.

Request deduplication remembers work already in progress.

Traffic distribution remembers server health.

Connection pools remember open connections.

Every optimization depends on memory.

And memory has one unavoidable weakness.

It gets older.

Tomorrow...

we're going to discover what happens when the entire architecture starts making decisions based on observations that are already out of date.

Not because the software is buggy.

Because time itself has become part of the system.

ENGINEER'S NOTEBOOK

Today's incident didn't teach me about load balancers.

It taught me about time.

I used to believe distributed systems were difficult because computers were distributed.

*Today I realized they're difficult because **the present is distributed.***

Every machine is living in a slightly different version of "now."

CHAPTER 12

The System Starts Watching Itself

For months, every architectural decision came from an engineer.

You added Redis.

You introduced another application server.

You configured traffic distribution.

You tuned PostgreSQL.

The system never made decisions.

It simply followed yours.

That era is ending.

. . .

Traffic continues to grow.

Not explosively.

Predictably.

Every week looks remarkably similar.

Monday is quiet.

Tuesday is slightly busier.

Friday morning brings another examination result.

The graphs have become familiar.

Almost comforting.

Until one morning...

they don't.

• • •

At 9:06 AM, CPU usage rises above eighty percent.

Nobody panics.

It has happened before.

At 9:08 AM...

it reaches ninety percent.

Still manageable.

At 9:09 AM...

response times begin increasing.

Users notice.

At 9:10 AM...

CPU suddenly falls back to thirty percent.

Nobody touched anything.

No deployment.

No restart.

No incident response.

The system recovered by itself.

. . .

The operations team is confused.

"What happened?"

Someone finally notices.

A third application server appeared.

Quietly.

Automatically.

Without anyone logging into a machine.

Without opening a terminal.

Without asking permission.

. . .

Yesterday, adding capacity was a decision.

Today...

it's a reaction.

The architecture has crossed another invisible line.

It is no longer waiting for engineers to tell it what to do.

It has started observing itself.

. . .

That sounds almost magical.

It isn't.

Imagine driving a car uphill.

As the hill becomes steeper...

you press the accelerator.

The car speeds up.

You don't calculate the exact engine torque required.

You observe.

You react.

Feedback.

Distributed systems eventually learn the same habit.

Observe.

React.

Observe again.

Adjust.

Repeat.

. . .

Notice something subtle.

The new application server didn't appear because traffic increased.

Traffic had been increasing for weeks.

It appeared because the system believed more capacity was necessary.

That's a different statement.

One describes reality.

The other describes an opinion about reality.

And opinions...

can be wrong.

Automation is simply architecture acting on its own beliefs.

. . .

The next morning, another traffic spike arrives.

CPU rises.

The system watches.

A fourth application server appears.

Traffic settles.

Everything looks perfect.

Until noon.

Traffic drops.

The servers remain.

Four machines.

Almost idle.

Someone asks,

"Why didn't they disappear?"

A few minutes later...

they do.

The system watched itself again.

Reached a different conclusion.

Acted accordingly.

For the first time...

your architecture is continuously changing shape.

. . .

That realization is exhilarating.

And terrifying.

Yesterday, every architectural change appeared in a pull request.

Today...

changes happen while you're drinking coffee.

Nobody reviewed them.

Nobody approved them.

Nobody even noticed.

Until they were already part of production.

. . .

The afternoon incident is surprisingly small.

One application server begins consuming excessive memory.

Normally, an engineer would investigate.

Collect heap dumps.

Analyze garbage collection.

Restart the service.

Instead...

the system quietly removes it from service.

Launches another.

Traffic continues.

Users never know.

The incident never reaches Slack.

Success has become invisible.

. . .

Then someone asks a question that changes the conversation.

"How does the system know when it's wrong?"

Silence.

Because nobody has considered it.

If CPU briefly spikes because of garbage collection...

should another server appear?

If network latency rises for ten seconds...

should traffic move elsewhere?

If one machine becomes temporarily slow...

should it be replaced?

The architecture is making decisions.

But every decision depends on observations.

And observations can be misleading.

. . .

You remember yesterday's lesson.

Information ages.

Today's lesson is different.

Interpretation fails.

The system doesn't react to reality.

It reacts to measurements.

Measurements are merely shadows cast by reality.

Sometimes...

shadows deceive us.

. . .

That evening, you review the day's graphs.

Every action taken by the architecture began with one number.

CPU.

Latency.

Queue length.

Memory.

Request rate.

The architecture never saw users.

It never saw code.

It never saw the incident.

It only saw measurements.

Then it acted.

You close your notebook and write one sentence.

"Every automated system is only as intelligent as the signals it believes."

Tomorrow...

those signals will begin disagreeing with one another.

One graph will say the system is healthy.

Another will insist it's failing.

Both will be correct.

And that contradiction will introduce us to one of the most misunderstood disciplines in software engineering.

Observability.

ENGINEER'S NOTEBOOK

I used to believe automation removed humans from the loop.

Today I realized it simply moved our decisions into software.

Every threshold...

Every timeout...

Every scaling policy...

Every retry budget...

*is a human opinion that continues making decisions long after
we've gone home.*

We don't remove ourselves from distributed systems.

We leave pieces of ourselves behind.

CHAPTER 13

Four Witnesses, One Incident

The alert arrives at 2:17 AM.

High latency.

No explanation.

No context.

Just a number crossing a threshold.

You open your laptop.

The architecture looks familiar.

Application servers.

Traffic distribution.

Redis.

PostgreSQL.

Nothing unusual.

Every dashboard is green.

Except one.

Latency.

. . .

You zoom into the graph.

Something definitely happened.

Response time climbed from 120 milliseconds to nearly three seconds.

Then...

without intervention...

it slowly recovered.

The graph answers one question beautifully.

Something changed.

It answers another question terribly.

Why?

. . .

Someone opens the application logs.

Thousands of lines begin flowing across the screen.

Requests.

Responses.

Authentication.

Database queries.

Warnings.

Retries.

Errors.

The information is there.

Somewhere.

Finding it feels like searching for a single conversation inside a crowded football stadium.

The logs answer a different question.

What happened?

Not necessarily...

What mattered?

. . .

Another engineer checks resource metrics.

CPU looks normal.

Memory stable.

Disk healthy.

Network traffic unchanged.

Interesting.

If the infrastructure wasn't struggling...

why were the users waiting?

The metrics tell a convincing story.

The application tells another.

Both appear true.

. . .

A fourth engineer opens a request trace.

One request.

Beginning to end.

The browser sends a request.

Traffic distribution forwards it.

Application A processes it.

Redis returns immediately.

PostgreSQL responds quickly.

Then...

almost two and a half seconds disappear.

Nothing happens.

The request simply waits.

Eventually it continues.

Eventually it finishes.

The trace reveals something neither graphs nor logs could show.

Not just **what** happened.

But **where time disappeared**.

. . .

By now the team has assembled four different explanations.

The graph says latency increased.

The logs say requests completed successfully.

The metrics say the servers were healthy.

The trace says the request spent most of its life waiting.

None of them disagree.

None of them completely agree.

Each is describing the same incident from a different perspective.

Like four witnesses describing the same car accident.

. . .

Imagine asking four people what happened at an intersection.

One tells you traffic became slower.

Another remembers hearing tires screech.

A third noticed the traffic lights never failed.

The fourth watched one particular car from beginning to end.

None of them are wrong.

None saw everything.

Production incidents work the same way.

Every observability signal is a witness, not the truth.

. . .

That realization changes the investigation.

You stop searching for the "correct" dashboard.

Instead...

you ask a different question.

"What kind of witness am I interviewing?"

Graphs summarize.

Logs remember.

Metrics measure.

Traces reconstruct journeys.

Each sacrifices something to become useful.

A graph sacrifices detail.

A log sacrifices simplicity.

A metric sacrifices context.

A trace sacrifices breadth.

Every perspective gains clarity by giving something else away.

. . .

An hour later, someone notices a pattern.

Every slow request shares one characteristic.

They all waited for the same external payment service.

Not because the payment service failed.

Because it became slightly slower than usual.

Each individual request only waited a little longer.

Thousands of simultaneous requests...

amplified that delay into a visible incident.

The application wasn't unhealthy.

It was patient.

Users weren't.

. . .

You resolve the incident.

The dashboards slowly return to normal.

The graphs flatten.

The logs become quiet.

The traces shorten.

Another peaceful night returns.

Before closing your laptop, you pause.

For the first time, you realize something that no architecture diagram has ever shown.

Your production system doesn't consist only of servers and databases.

It also consists of the stories those machines tell about themselves.

Those stories are incomplete.

Sometimes contradictory.

Always valuable.

Tomorrow...

we're going to discover what happens when two perfectly healthy services begin waiting for each other.

Not because either one is broken.

Because neither one knows when to stop waiting.

ENGINEER'S NOTEBOOK

Tonight I stopped looking for "the dashboard."

There isn't one.

Every production incident is a conversation between partial truths.

*Metrics tell me **that** something changed.*

*Logs tell me **what** happened.*

*Traces tell me **where** time went.*

None of them tell the whole story.

Neither do engineers.

CHAPTER 14

The Waiting Game

The alert doesn't look urgent.

Request latency has increased again.

Nothing unusual.

It has happened before.

You open the dashboards expecting another familiar story.

High CPU.

Busy database.

Network congestion.

Instead...

everything looks healthy.

Application servers are comfortably below forty percent CPU.

Redis responds almost instantly.

PostgreSQL queries remain fast.

Memory usage is stable.

The network shows no signs of congestion.

Every component appears to be doing exactly what it should.

Users, however, disagree.

. . .

Support tickets begin arriving.

"The page keeps loading."

"It never finishes."

"I tried refreshing."

Nothing crashes.

Nothing fails.

Nothing completes.

. . .

You pick one request and follow its journey.

The browser sends the request.

Traffic distribution forwards it.

The application begins processing.

It asks another internal service for additional information.

Then...

it waits.

Not forever.

Just a little longer than usual.

The downstream service eventually replies.

The request continues.

Nothing remarkable.

. . .

You inspect another slow request.

The pattern repeats.

The application asks another service.

Then waits.

A third request.

Waits.

A fourth.

Waits.

Soon, every slow request shares the same characteristic.

Not heavy computation.

Waiting.

. . .

Someone suggests increasing CPU.

Another proposes larger machines.

Neither idea explains what you're seeing.

The processors aren't busy.

They're idle.

The application isn't working.

It's waiting for permission to continue.

. . .

Imagine two people trying to leave a narrow hallway.

Each steps aside to let the other pass.

Each waits politely.

Neither moves.

Nothing blocks the hallway.

Courtesy does.

Distributed systems often fail with the same elegance.

Waiting is work that nobody notices.

. . .

The requests continue accumulating.

Not because processing became slower.

Because every request spends more time existing.

A request that finishes in one hundred milliseconds occupies almost nothing.

A request that waits five seconds occupies memory.

Connections.

Threads.

Attention.

Waiting quietly consumes resources without performing useful work.

. . .

Eventually...

something surprising happens.

The database connection pool becomes exhausted.

Not because PostgreSQL slowed down.

Because requests held their connections while waiting for another service.

Soon...

new requests cannot obtain a database connection.

Those requests begin waiting too.

The waiting spreads.

Like traffic slowing several kilometers behind an accident.

Most of the cars were never near the collision.

They merely inherited its consequences.

. . .

The architecture diagram hasn't changed.

No new components.

No new deployments.

No hardware failures.

Only one thing changed.

The amount of time each service spent waiting for another.

Small delays became larger queues.

Larger queues created longer waits.

Longer waits created even larger queues.

The architecture quietly began amplifying itself.

. . .

This is the first time you've encountered a feedback loop that hurts instead of helps.

Yesterday...

feedback created new application servers.

Today...

feedback creates congestion.

The system isn't failing because individual components are unhealthy.

It's failing because healthy components are influencing one another.

The architecture has become connected strongly enough for local problems to become global ones.

. . .

An engineer asks a question that changes the investigation.

"How long should we wait before giving up?"

Silence.

Nobody has ever written down the answer.

The applications have always waited.

Patiently.

Politely.

Forever, if necessary.

For the first time, infinite patience looks irresponsible.

. . .

You configure a timeout.

Not because you expect failures.

Because you finally accept that waiting forever is itself a failure.

The next incident looks different.

Some requests now fail quickly.

Users occasionally retry.

Most requests succeed.

The architecture remains responsive.

The system has learned something profoundly human.

Sometimes saying "**I can't**" is healthier than endlessly saying "**Maybe.**"

. . .

The incident closes.

Not because every request succeeded.

Because the architecture stopped allowing one slow service to immobilize everything else.

As you pack your laptop into your bag, one realization refuses to leave.

The hardest part of building distributed systems isn't teaching machines how to communicate.

It's teaching them when to stop waiting for each other.

Tomorrow...

we're going to discover what happens when stopping too early becomes just as dangerous as waiting too long.

Because every timeout eventually becomes another engineer's retry.

Today I learned that a system can be perfectly healthy and completely unavailable.

Health belongs to components.

Availability belongs to relationships.

A service can spend its entire life waiting politely for another service.

Enough polite services...

and an entire architecture grinds to a halt.

CHAPTER 15

Hope Is Not a Strategy

The timeout works.

Exactly as designed.

Instead of waiting forever, the application gives up after three seconds.

Users occasionally see an error.

Most don't.

The architecture remains responsive.

The incident ends faster than the last one.

Everyone sleeps better.

Until the next morning.

. . .

The support team forwards a familiar message.

"Sometimes refreshing fixes it."

You smile.

Of course it does.

The first request timed out.

The second one happened to arrive after the downstream service recovered.

Nothing mysterious.

Just better timing.

Someone around the table asks the obvious question.

"If trying again often works...

why don't we let the software try again automatically?"

The room nods.

It sounds reasonable.

Helpful, even.

. . .

The change is surprisingly small.

When a request fails...

try once more.

No complicated architecture.

No new infrastructure.

Just another attempt.

The deployment goes live.

The graphs look unchanged.

Success.

. . .

For three weeks, nobody thinks about retries.

Failures quietly recover.

Transient network hiccups disappear.

Brief service restarts become invisible.

The architecture feels more reliable.

Not because failures stopped happening.

Because users stopped noticing them.

Retries have quietly become heroes.

. . .

Then Friday arrives.

At exactly 10:01 AM, a downstream service begins responding slowly.

Not failing.

Just slower than usual.

The first requests time out.

The application politely tries again.

The second attempt also waits.

More requests arrive.

Each one now becomes two.

Traffic hasn't doubled.

Attempts have.

. . .

The downstream service was already struggling.

Now every delayed request creates another request.

The service receives more work precisely when it has become less capable of doing work.

The architecture begins feeding its own problem.

. . .

Imagine a crowded restaurant.

Meals begin taking longer.

Customers grow impatient.

Instead of waiting...

everyone places the same order again.

The kitchen doesn't receive more hungry people.

It receives more tickets.

The cooks don't know which order is the latest.

They simply see more work.

The delays become even longer.

More duplicate orders arrive.

Soon the restaurant isn't overwhelmed by customers.

It's overwhelmed by hope.

Every retry assumes the next attempt will have a better future than the last.

Sometimes that's true.

Sometimes the future is simply more crowded.

. . .

Back in production, dashboards become confusing.

Request volume doubles.

Unique users remain unchanged.

CPU rises.

Network traffic increases.

The downstream service begins timing out even more frequently.

Each timeout creates another retry.

Each retry creates another timeout.

Nobody introduced a bug.

Nobody deployed broken code.

The architecture simply became too optimistic.

. . .

The operations team disables retries.

Almost immediately...

traffic falls.

The downstream service slowly catches its breath.

Requests begin succeeding again.

Ironically...

the quickest way to recover was to stop trying so hard.

. . .

That afternoon, someone sketches two arrows on a whiteboard.

Failures.

Retries.

Then another arrow.

Retries.

Failures.

The arrows form a circle.

Nobody writes the words "feedback loop."

Nobody needs to.

Everyone can already see it.

. . .

The lesson lingers long after the incident ends.

Retrying wasn't wrong.

Retrying without understanding *why* something failed was.

A network packet lost during transmission deserves another chance.

A service drowning under heavy load probably doesn't.

Hope is valuable.

Blind optimism is expensive.

. . .

As you leave the office, one thought refuses to disappear.

Every retry asks the same question.

"Has the world changed enough for this attempt to succeed?"

Sometimes the answer is yes.

Sometimes it's no.

Tomorrow...

we're going to discover how software learns to stop asking that question altogether.

Because eventually...

the healthiest decision isn't to retry.

It's to stop calling.

Today I realized that retries don't create resilience.

They create another conversation with the future.

Sometimes the future has recovered.

Sometimes it has become worse.

Software shouldn't retry because it is hopeful.

It should retry because it has evidence that trying again makes sense.

CHAPTER 16

Learning to Say No

For weeks, the retry policy remained disabled.

The architecture became quieter.

Request volume matched user volume again.

Incidents became shorter.

The graphs felt predictable.

Predictability has a dangerous side effect.

It convinces you the problem is solved.

. . .

Monday morning arrives with another examination result.

Traffic rises exactly as expected.

Redis performs well.

PostgreSQL remains comfortable.

Application servers scale automatically.

Everything you've built over the past months works together.

Then...

the payment service begins slowing down.

Again.

Not failing.

Not crashing.

Just answering more slowly than usual.

. . .

The application sends a request.

It waits.

Three seconds pass.

The timeout expires.

This time...

it doesn't retry.

Exactly as planned.

The request fails.

The application immediately accepts another incoming request.

Then another.

Then another.

Every one of them asks the same downstream service.

Every one of them eventually fails.

The architecture has stopped repeating mistakes.

It hasn't stopped making them.

. . .

Support tickets begin arriving.

Users don't understand why one particular feature no longer works.

Everything else feels normal.

Login succeeds.

Search works.

Results load.

Only payments fail.

. . .

An engineer suggests something unexpected.

"What if we stop calling the payment service?"

The room looks puzzled.

"How would payments ever recover if we never ask again?"

It's a fair question.

. . .

Imagine calling a friend.

They don't answer.

You call again.

Still nothing.

Again.

Again.

Eventually...

the calls themselves become the problem.

Your friend cannot answer because the phone never stops ringing.

Sometimes the kindest thing you can do is stop calling for a while.

Distributed systems eventually learn the same lesson.

Mercy is sometimes more useful than persistence.

. . .

The change is surprisingly simple.

After enough consecutive failures...

stop sending requests.

Not forever.

Just long enough for the downstream service to recover.

New requests fail immediately.

No waiting.

No retries.

No unnecessary work.

At first, it feels wrong.

The application is rejecting requests without even trying.

Engineers spend their careers making software do more.

This is the first feature whose purpose is to do less.

. . .

The deployment reaches production.

The payment service slows again.

The first few requests behave exactly as before.

Timeout.

Failure.

Another timeout.

Another failure.

Then something changes.

New requests are rejected immediately.

The application doesn't wait.

The users don't stare at loading indicators.

The downstream service finally receives something it hasn't seen all morning.

Silence.

. . .

For thirty seconds...

almost nothing happens.

The payment service catches up.

Its queues begin shrinking.

CPU settles.

Response times recover.

The architecture has done something remarkable.

It improved another system simply by refusing to participate in its struggle.

. . .

One engineer frowns at the dashboard.

"We're intentionally returning errors."

Yes.

Another engineer replies quietly,

"But they're fast errors."

The room goes silent.

Because everyone suddenly understands the trade.

A five-second failure feels like uncertainty.

A fifty-millisecond failure feels like honesty.

Users rarely enjoy errors.

They almost always prefer certainty.

. . .

A minute later, the application cautiously sends another request.

It succeeds.

Another succeeds.

Then another.

Traffic gradually resumes.

Nothing dramatic.

No celebration.

Just patience.

The architecture learned to ask a different question.

Not,

"Can the payment service answer?"

But,

"Is it fair to ask right now?"

. . .

That distinction changes the way you think about distributed systems.

Until now, every optimization tried to maximize success.

Now you're deliberately limiting work.

Not because you're pessimistic.

Because survival sometimes requires restraint.

. . .

As the incident closes, you realize something unexpected.

The architecture has become capable of protecting itself.

Not by becoming stronger.

By recognizing when strength is no longer the right strategy.

Tomorrow...

we're going to discover that saying "no" isn't enough.

Because once dozens of independent services begin protecting themselves...

they also have to agree on **who should act first**.

The next challenge won't be failure.

It will be leadership.

ENGINEER'S NOTEBOOK

Today I stopped thinking about failures.

I started thinking about boundaries.

Good engineers know how to build powerful systems.

Great engineers know where those systems should stop.

Every "no" in a distributed system protects a future "yes."

CHAPTER 17

When Everyone Decides

The payment incident fades into memory.

The architecture has become calmer.

Healthier.

More disciplined.

Services know when to wait.

When to stop waiting.

When to retry.

When to refuse.

Months earlier, every decision belonged to an engineer.

Now...

many decisions belong to the system itself.

That feels like progress.

Until one Tuesday morning.

. . .

At exactly 11:00 AM, a scheduled task begins.

Its job is simple.

Generate examination rankings.

Millions of records.

One report.

Nothing interactive.

No users waiting.

Just background work.

For years, a single application server handled the task.

Now there are six.

Every server starts at the same time.

Every server notices the schedule.

Every server reaches the same conclusion.

"It's my responsibility."

. . .

Nobody programmed them to compete.

Nobody programmed them to cooperate either.

Each server simply follows the same logic.

The clock reaches eleven.

The job begins.

Six times.

. . .

At first, nothing seems unusual.

CPU increases.

Database activity rises.

Storage writes accelerate.

Then someone notices something strange.

The ranking table is growing far faster than expected.

Duplicates appear.

Reports overwrite one another.

Logs from different servers proudly announce,

"Ranking generation completed successfully."

Every server believes it succeeded.

They're all telling the truth.

The system, however, is wrong.

. . .

The incident isn't caused by failure.

It's caused by agreement.

Every server made the same perfectly reasonable decision.

The architecture didn't need better judgment.

It needed fewer judges.

Coordination doesn't begin when systems disagree. It begins when they agree too often.

. . .

Someone asks the obvious question.

"Why don't we assign one server to perform the job?"

Reasonable.

Until that server crashes.

Then who takes over?

Another engineer suggests assigning a backup.

Then a backup for the backup.

The whiteboard slowly fills with arrows.

Each new solution quietly creates another responsibility.

Who decides when the backup becomes primary?

. . .

You realize the architecture has reached another turning point.

Earlier chapters asked,

"Who should answer this request?"

Today's question is different.

"Who should make this decision?"

Those sound similar.

They aren't.

Requests happen millions of times every day.

Leadership changes rarely.

One distributes work.

The other distributes responsibility.

. . .

Imagine six firefighters arriving at the same emergency.

If every firefighter independently decides to command the operation...

chaos follows.

If none of them does...

chaos follows.

The problem isn't competence.

The problem is clarity.

Distributed systems eventually discover the same truth.

Capability without coordination creates conflict.

. . .

The team introduces one simple rule.

Before beginning the scheduled task...

ask whether someone else is already responsible.

If the answer is yes...

stand down.

If the answer is no...

accept responsibility.

The rule feels almost disappointingly ordinary.

Until the next deployment.

. . .

One server begins generating rankings.

The others quietly wait.

Halfway through the job...

the responsible server crashes.

Silence.

For a brief moment...

nobody is responsible.

Then another server notices.

It cautiously accepts the responsibility.

The work resumes.

No duplicates.

No confusion.

Just continuity.

. . .

The architecture has learned something profound.

Responsibility is temporary.

It belongs to whoever can fulfill it today.

Not whoever fulfilled it yesterday.

Leadership isn't ownership.

It's stewardship.

. . .

As the afternoon incident review concludes, someone points toward the architecture diagram.

Nothing has changed.

The same servers.

The same databases.

The same caches.

Yet the architecture behaves completely differently than it did yesterday.

Not because new software was written.

Because every server finally understands something that humans have always known.

Sometimes...

the most important decision isn't **what** to do.

It's deciding **who should do it.**

Tomorrow...

we're going to discover something even stranger.

Sometimes the system can't even agree on who is alive.

And that uncertainty will force us to confront one of the deepest questions in distributed computing.

Not leadership.

Truth.

ENGINEER'S NOTEBOOK

Today I learned that responsibility is another finite resource.

Giving everyone authority feels democratic.

It also creates confusion.

Distributed systems don't fail because nobody can lead.

They fail because too many machines sincerely believe they already are.

CHAPTER 18

The Day Truth Split in Two

At 2:43 PM, Server Three disappears.

Not dramatically.

No alarms.

No smoke.

No crash report.

It simply stops answering.

. . .

The remaining servers continue working.

Requests still succeed.

Background jobs continue.

Users notice nothing unusual.

Yet inside the architecture...
an argument has already begun.

. . .

Server One believes Server Three has failed.
It hasn't received a heartbeat for several seconds.
Server Two isn't so sure.
Its last heartbeat arrived only moments ago.
Perhaps the network is slow.
Perhaps the messages are delayed.
Perhaps Server Three is still perfectly healthy.
Two intelligent machines.
Two different conclusions.
Both based on honest observations.

. . .

An engineer joins the incident bridge.
"Is Server Three down?"
Nobody answers immediately.

Not because they don't know how to investigate.

Because the question itself is harder than it sounds.

. . .

Imagine standing on opposite sides of a mountain.

You watch the sun disappear.

Someone on the other side still sees daylight.

Neither of you is lying.

Neither of you is mistaken.

You simply occupy different positions.

Distributed systems live on opposite sides of that mountain every second of every day.

There is no single viewpoint from which the entire system can be observed.

. . .

The team restarts Server Three.

Nothing changes.

Server One still ignores it.

Server Two begins sending work to it again.

The architecture has quietly divided into two different realities.

Neither reality is complete.

. . .

Someone opens the monitoring dashboard.

It confidently reports:

Five healthy servers.

Someone else checks the deployment dashboard.

It reports:

Four active servers.

Both dashboards were built correctly.

Both queried different sources.

Both tell different stories.

Which one should the team believe?

. . .

Until today, you've trusted every graph.

Every metric.

Every alert.

Now you realize something uncomfortable.

Every measurement has an observer.

Every observer has limitations.

Truth depends on perspective.

. . .

The discussion changes.

Nobody asks,

"Which dashboard is correct?"

Instead they ask,

"Who observed this?"

That tiny change transforms the investigation.

Evidence matters.

So does the path by which evidence arrived.

. . .

Eventually the networking team discovers the problem.

A switch between two racks had briefly failed.

Server Three never stopped running.

It merely became unreachable from part of the architecture.

Half the system believed it had disappeared.

The other half continued talking to it normally.

The server wasn't dead.

Communication was.

. . .

The incident ends.

No data was lost.

No users noticed.

Still...

everyone leaves the meeting unsettled.

For years, the team assumed failures were obvious.

Servers crashed.

Processes exited.

Disks filled.

Today's failure was different.

The server and the network each told a believable story.

Those stories simply couldn't both describe the same reality.

. . .

You begin reviewing previous incidents.

Suddenly...

old mysteries make more sense.

The stale cache.

The delayed health check.

The overloaded application server.

The duplicated background job.

Every one of them shared the same hidden theme.

Different parts of the architecture believed different versions of the present.

This wasn't a new problem.

It was the first time you recognized it.

. . .

Late that evening, someone asks the question nobody wanted to ask.

"If intelligent machines observing the same system can disagree...
who gets the final say?"

The room grows quiet.

Because nobody wants one server deciding the truth for everyone else.

Neither do they want every server deciding independently.

Somehow...

the architecture needs a way to transform many uncertain opinions into one dependable decision.

Tomorrow...

we're going to discover why agreement is far more difficult than correctness.

Because in distributed systems...

being right is often easier than convincing everyone else that you are.

Today I stopped asking whether the system was healthy.

*I started asking **who believed it was healthy.***

Every dashboard is someone's point of view.

Every alert is an opinion formed from incomplete evidence.

Truth doesn't disappear in distributed systems.

It fragments.

CHAPTER 19

When Nobody Gets to Decide Alone

The incident is over.

Server Three never failed.

The network simply stopped carrying its voice.

By the afternoon, every server agrees again.

Traffic flows normally.

Users remain unaware.

The architecture appears peaceful.

Appearances can be misleading.

. . .

The incident review begins with a simple proposal.

"If one server isn't enough to decide the truth...

why don't we let all the servers vote?"

Heads nod around the room.

It sounds fair.

Democratic.

Reasonable.

Surely five intelligent machines can reach a better decision than one.

. . .

The next week brings another network interruption.

Smaller this time.

Server One cannot communicate with Server Four.

Server Four cannot communicate with Server One.

The remaining three servers communicate normally.

Each machine asks the same question.

"Am I alone?"

Each receives a different answer.

. . .

Server One counts four reachable machines.

Server Four also counts four.

Both conclude that they belong to the healthy majority.

Both continue accepting responsibility.

Neither realizes the other has reached the same conclusion.

. . .

Nothing crashes.

Nothing restarts.

Background jobs continue.

Requests succeed.

Then...

someone notices two different reports for the same examination.

Generated only seconds apart.

Both marked as official.

Both internally consistent.

Both different.

. . .

The architecture didn't produce an incorrect answer.

It produced two correct answers.

Each was correct according to the world its author could observe.

***Truth without agreement is indistinguishable
from contradiction.***

. . .

Imagine two newspapers printed in different cities.

A bridge between the cities collapses overnight.

Each newspaper reports the events it witnessed.

Both editions are accurate.

Neither edition is complete.

When the bridge reopens...

which newspaper becomes history?

Distributed systems ask this question every day.

. . .

The engineers gather around the whiteboard.

Someone circles a single sentence.

"There must never be two official answers."

The room becomes quiet.

Not because the sentence is difficult.

Because everyone suddenly understands its cost.

If two parts of the architecture cannot communicate...

someone may have to stop making decisions.

Not because they're broken.

Because certainty has become impossible.

. . .

The idea feels unfair.

A perfectly healthy server may deliberately refuse useful work.

A database with no errors may reject writes.

An application with plenty of CPU may remain idle.

Not because it lacks capability.

Because capability is no longer the most important requirement.

Agreement is.

. . .

The architecture changes.

Before accepting responsibility...

a server first asks,

"Can enough of us still see each other?"

If the answer is uncertain...

it chooses silence.

The decision feels wasteful.

Until the next incident.

. . .

Another network interruption.

The same separation.

The same uncertainty.

This time...

one side quietly stops making changes.

Users receive temporary errors.

Background work pauses.

The other side continues.

No conflicting reports.

No duplicated decisions.

No competing truths.

The architecture chose consistency over activity.

For the first time...

doing less protected correctness.

. . .

The incident ends.

Communication returns.

The paused servers rejoin.

Nothing dramatic happens.

No complicated recovery.

Just quiet agreement.

As the systems settle back into rhythm, you realize something unexpected.

The architecture never learned how to become smarter.

It learned how to become humbler.

Sometimes...

the most trustworthy machine in the room is the one willing to admit,

"I don't know enough to decide."

. . .

You close the incident review.

One sentence remains on the whiteboard.

"No decision is better than two conflicting ones."

It feels strangely unsatisfying.

Users experienced errors.

Requests were rejected.

Work was delayed.

Yet everyone leaves believing the system behaved correctly.

Tomorrow...

we're going to discover why protecting one truth is only half the story.

Because once everyone finally agrees...

someone has to remember that agreement long after today's machines have disappeared.

Today I stopped thinking of consensus as voting.

Voting counts opinions.

Consensus protects reality.

*Sometimes the strongest decision a distributed system can make
is refusing to pretend it knows something it cannot possibly
know.*

Confidence without evidence is just another failure mode.

CHAPTER 20

The Memory Everyone Must Share

The network heals.

The servers agree again.

Background jobs resume.

Users stop seeing errors.

Another incident quietly becomes another story.

The architecture feels older now.

Not larger.

Wiser.

It has learned to question certainty before acting.

That lesson comes with an unexpected consequence.

• • •

During the incident review, someone asks a simple question.

"How do we know everyone agrees on the same report tomorrow?"

The room pauses.

Today's agreement feels reassuring.

Tomorrow is another day.

Servers restart.

Deployments happen.

Machines disappear.

New ones appear.

Agreement is temporary.

Memory must survive it.

. . .

Imagine finishing a meeting where everyone finally agrees on an important decision.

Then everyone leaves the room.

The next morning...

someone asks,

"What did we decide yesterday?"

If nobody wrote it down...

agreement lasted exactly as long as the meeting.

Distributed systems eventually discover the same weakness.

Agreement without memory is merely conversation.

Truth survives only if it is remembered.

. . .

For years, writing to the database felt ordinary.

An UPDATE statement.

A transaction.

A commit.

The request returned.

Life moved on.

Today...

the team looks at a successful write differently.

A successful write is no longer just new information.

It is a promise to the future.

. . .

Every report generated.

Every payment completed.

Every examination result published.

Every inventory adjustment.

Someone, somewhere, will make another decision based on that write.

Tomorrow.

Next week.

Years later.

The write isn't serving today's request.

It's shaping tomorrow's reality.

. . .

An engineer sketches two words on the whiteboard.

Read.

Write.

Then draws a thick line beneath one of them.

Write.

The room smiles.

Everyone understands why.

Reading can happen endlessly.

A million times.

A hundred million times.

Each reader merely observes.

Writing changes the story everyone else will read.

The architecture doesn't fear readers.

It fears conflicting authors.

. . .

You begin reviewing the previous months.

The stale cache.

The duplicated report.

The competing leaders.

The partitioned network.

Every incident threatened the same thing.

Not availability.

Not performance.

Shared memory.

The architecture wasn't protecting CPUs.

It was protecting history.

. . .

That realization quietly changes every future design discussion.

The first question is no longer,

"How fast can we read this?"

Instead it becomes,

"Who is allowed to change it?"

Those questions produce entirely different architectures.

. . .

The next feature request arrives.

Publish updated examination rankings.

Simple enough.

Someone asks,

"What if two administrators submit different corrections at the same time?"

Nobody reaches for another server.

Nobody asks about CPU.

Nobody suggests more Redis.

Everyone asks the same question.

"Which write becomes history?"

The discussion has changed forever.

. . .

Late that afternoon, the first correction is accepted.

The second arrives milliseconds later.

The system pauses.

Not because it is slow.

Because changing shared memory deserves hesitation.

A hundred milliseconds spent protecting one write can prevent years of confusion.

The architecture has become careful.

Not cautious.

Careful.

There is a difference.

. . .

You leave the office looking at the title on your notebook.

One Write.

One Hundred Million Reads.

For months, you believed the difficult part was serving millions of readers.

Today...

you understand something different.

Readers are easy.

Readers only ask.

Writers decide what everyone else will believe.

Tomorrow...

we're going to discover why protecting a single write sometimes requires thousands of machines to briefly stop trusting themselves.

Because preserving one shared memory is harder than creating a thousand private ones.

ENGINEER'S NOTEBOOK

Today I stopped treating writes as database operations.

They are promises.

Every successful write tells the future,

"This is the story we have all agreed to remember."

Reads consume history.

Writes create it.

CHAPTER 21

The Half-Written Story

The correction request arrives just before lunch.

An administrator discovers a scoring mistake.

One candidate's result must change.

The operation is routine.

The application validates the request.

Updates the examination record.

Prepares to notify downstream systems.

Everything proceeds exactly as planned.

Until it doesn't.

. . .

Halfway through the operation...

the application server crashes.

No warning.

No graceful shutdown.

Just silence.

. . .

A few seconds later, another application server replaces it.

Users continue browsing.

Search still works.

Reports still load.

From the outside...

production appears healthy.

Inside the architecture...

a single question begins echoing.

What actually happened?

. . .

Did the correction reach the database?

Did the notification get sent?

Did other services already begin updating caches?

Did the administrator receive success?

Did the browser retry?

Nobody knows.

. . .

Someone opens PostgreSQL.

The updated score exists.

Another engineer checks the notification service.

No event was received.

Redis still contains the previous value.

The search index hasn't changed.

Some parts of the architecture remember yesterday.

Others remember today.

The write succeeded.

The story didn't.

. . .

For months, every chapter taught you something different.

Resources.

Coordination.

Leadership.

Truth.

Memory.

Today...

all of them collide.

. . .

Imagine an author writing the final chapter of a novel.

Halfway through the last page...

the lights go out.

The manuscript survives.

But only part of the final chapter reaches the publisher.

Readers receive different endings.

Nobody argues about the words that exist.

They argue about the ones that never arrived.

Distributed systems don't fear failed writes.

They fear incomplete history.

*The most dangerous write is the one that
succeeded only halfway.*

. . .

The incident room becomes unusually quiet.

Nobody asks how to restart the server.

Nobody asks about CPU.

Nobody asks about memory.

Those questions suddenly feel irrelevant.

The architecture has encountered something much more
unsettling.

It no longer knows where one story ends and another begins.

. . .

An engineer asks,

"Can we simply perform the write again?"

Maybe.

Unless the first write already happened.

Another engineer asks,

"Can we undo everything?"

Maybe.

Unless another service has already acted on the first version.

Every proposed solution depends on answering the same impossible question.

"What is the current truth?"

. . .

The architecture has spent months learning how to agree.

Now agreement isn't enough.

Agreement must survive interruption.

A write cannot merely become correct.

It must become undeniable.

. . .

That realization changes every future design discussion.

Nobody talks about performance.

Nobody mentions latency.

Every conversation circles the same idea.

If tomorrow's machines know nothing about today's failure...
how will they continue today's story?

. . .

You notice something remarkable.
Every great distributed system eventually invents the same idea.
Not because engineers copy each other.
Because reality leaves them little choice.
The architecture needs a memory that survives individual machines.
A memory that cannot disappear when a process exits.
A memory that tells every future server exactly where yesterday ended.
Not because remembering is elegant.
Because forgetting is catastrophic.

. . .

The incident closes hours later.
Not because engineers became faster.
Because they finally reconstructed the missing history.

One write.

Hundreds of log files.

Thousands of timestamps.

Dozens of machines.

The irony isn't lost on anyone.

Creating history took milliseconds.

Recovering it took an afternoon.

. . .

You shut your laptop.

One sentence stays behind.

"If we cannot continue the story after failure...

...then we never truly finished writing it."

Tomorrow...

we're going to discover why the most important data structure in distributed systems isn't a database.

It's a diary.

A diary every machine trusts more than its own memory.

ENGINEER'S NOTEBOOK

Today I learned that persistence is not about storing data.

It is about preserving unfinished conversations.

Machines fail.

Processes disappear.

Engineers go home.

History cannot.

*Every durable system is, at its heart, a promise that tomorrow
can continue yesterday's story.*

CHAPTER 22

The Diary

The incident is over.

The correction was eventually completed.

Caches were rebuilt.

Notifications were replayed.

Search indexes caught up.

Users never learned how close the architecture came to telling two different stories.

The engineers did.

. . .

The following morning, the incident review begins with a question no one has asked before.

"If this happened again tomorrow...

how would the new server know where yesterday stopped?"

Nobody answers immediately.

Because the question isn't about databases.

It's about memory.

. . .

Yesterday's application server is gone.

Its memory vanished the moment the process ended.

Its variables disappeared.

Its threads disappeared.

Its call stack disappeared.

Yet today's server is expected to continue yesterday's work as though nothing happened.

How?

. . .

Imagine inheriting a library after the librarian unexpectedly disappears.

Every book remains.

Every shelf remains.

But halfway through cataloguing a new collection, the librarian stopped.

No notes.

No checklist.

No journal.

Would you know which books had already been processed?

Would you know which ones were still waiting?

Eventually...

someone would begin keeping a notebook.

Not because notebooks are clever.

Because human memory is temporary.

Distributed systems discover the same necessity.

***Memory becomes trustworthy only after it leaves
memory.***

• • •

The engineers stop discussing machines.

Instead...

they begin discussing sequence.

What happened first?

What happened next?

Which decision came before another?

Which correction replaced the previous one?

Notice the change.

Nobody asks,

"What is the current value?"

They ask,

"How did we arrive here?"

History has quietly become more valuable than state.

. . .

The architecture changes in a surprisingly subtle way.

Instead of remembering only the latest answer...

it begins remembering the journey.

Every accepted correction.

Every completed payment.

Every published result.

Every decision.

Not because anyone expects to read them all.

Because tomorrow's machines may need to.

. . .

The first few days feel wasteful.

The diary keeps growing.

Millions of entries.

Most never examined again.

Someone asks,

"Why keep events we'll never read?"

The question feels reasonable.

Until another incident.

. . .

A deployment fails halfway through.

One application disappears.

Another immediately takes its place.

The new server doesn't ask,

"What was the old server thinking?"

It simply opens the diary.

Page by page.

Yesterday's decisions become today's understanding.

The work continues.

Not perfectly.

Not instantly.

But continuously.

The architecture no longer depends on remembering.

It depends on recording.

. . .

Weeks later, another administrator corrects an examination result.

This time, something feels different.

The correction isn't replacing history.

It's extending it.

Yesterday's result still exists.

So does today's.

One tells the story as it was.

The other tells the story as it became.

The architecture has stopped erasing its past.

It has started preserving it.

. . .

You begin noticing the same idea everywhere.

Account statements.

Medical records.

Legal contracts.

Airline reservations.

Financial transactions.

The most important systems in the world rarely destroy history.

They append to it.

Not because storage is cheap.

Because memory is expensive.

. . .

Late that evening, someone finally gives the notebook a name.

A log.

The word feels almost disappointingly ordinary.

Everyone nods.

Then returns to work.

Because the important idea was never the name.

The important idea is that tomorrow's machines no longer have to trust yesterday's machines.

They only have to trust what yesterday wrote down.

. . .

You close your own notebook before leaving.

For months, you've been writing observations after every incident.

Small reflections.

Questions.

Lessons.

They never felt particularly important.

Until today.

You smile.

Without realizing it...

you've been using the oldest distributed systems technique in history.

A diary.

Tomorrow...

we're going to discover why one diary is never enough.

Because memories can burn.

Buildings can flood.

Disks can fail.

And history that exists in only one place is still one failure away from disappearing forever.

ENGINEER'S NOTEBOOK

Today I stopped thinking of logs as debugging tools.

They are memory made durable.

Every machine eventually forgets.

Every process eventually ends.

Every engineer eventually hands the system to someone else.

A diary is how yesterday teaches tomorrow.

CHAPTER 23

The Cost of Forgetting

The diary changes everything.

Incidents become shorter.

New servers recover faster.

Deployments feel less frightening.

Yesterday finally has a way to speak to tomorrow.

For a while...

that feels like enough.

. . .

Until a hard drive fails.

No warning.

No gradual slowdown.

One moment the diary exists.

The next moment...

it doesn't.

. . .

The architecture continues running.

Users still browse results.

Search still works.

Payments continue processing.

At first glance...

nothing seems unusual.

Then someone asks a quiet question.

"Where is yesterday?"

. . .

The room falls silent.

Not because anyone doubts the answer.

Because everyone already knows it.

Yesterday lived inside the diary.

The diary lived on one machine.

The machine is gone.

. . .

Imagine discovering that the only handwritten constitution of a nation has been destroyed.

Every citizen remembers pieces of it.

Some remember the opening paragraphs.

Others recall important amendments.

Nobody remembers every word.

The nation still exists.

Its shared memory does not.

***History remembered by many minds is still
vulnerable if it was written only once.***

. . .

The engineers begin reconstructing events.

Application logs.

Database snapshots.

Monitoring dashboards.

Message queues.

Email confirmations.

Each contains fragments.

None contains the whole story.

By evening...

they've rebuilt most of yesterday.

Not all of it.

A handful of corrections remain uncertain.

Nobody can honestly say whether they happened.

Or almost happened.

. . .

The incident feels different from every previous one.

Nothing behaved incorrectly.

Nothing made a poor decision.

The architecture simply forgot.

And forgetting, they realize, is not an event.

It's permanent.

. . .

The next morning, someone places two notebooks on the conference table.

Identical.

"If this one disappears..."

they ask,

"...should history disappear with it?"

The answer arrives immediately.

"No."

Nobody argues.

Nobody hesitates.

The discussion isn't about performance.

It isn't about availability.

It's about memory deserving another chance.

. . .

The architecture changes again.

Not dramatically.

The diary remains.

A second copy quietly appears elsewhere.

Not for faster reads.

Not for backup reports.

Simply because one memory is no longer enough.

The second diary doesn't make history more important.

It makes history harder to erase.

. . .

Weeks later, another disk fails.

This time...

nobody panics.

The surviving diary opens.

The replacement machine quietly copies every page.

Yesterday returns.

Not because the failed machine recovered.

Because yesterday had already been entrusted to someone else.

. . .

The engineers notice something subtle.

Copying history isn't the difficult part.

Knowing **which** history to copy is.

The second diary must never invent its own version.

It must preserve the same story.

Character for character.

Decision for decision.

The architecture isn't protecting files.

It's protecting identity.

. . .

You leave the office thinking about families.

Parents tell stories to children.

Teachers preserve discoveries made centuries earlier.

Libraries outlive authors.

Civilizations survive because memory moves from one generation to the next.

Distributed systems quietly rediscovered an ancient truth.

Knowledge survives only after it has been shared.

. . .

That night, you write one final sentence before closing your notebook.

"History doesn't become durable when it's stored.

It becomes durable when it can survive being forgotten by one place."

Tomorrow...

we're going to discover why sharing history creates an entirely new problem.

When several diaries exist...

which one speaks for the future?

Because copying memory is easy.

Protecting one shared version of memory is where distributed systems earn their name.

ENGINEER'S NOTEBOOK

Today I realized that durability isn't about disks.

It's about inheritance.

One copy remembers.

Two copies preserve.

But preservation brings a new responsibility.

*Every copy must remember the **same** story.*

CHAPTER 24

One Write

The second diary works.

The failed machine returns.

History survives.

The architecture breathes a little easier.

For the first time, forgetting no longer feels inevitable.

Then another question quietly appears.

One nobody asked before.

. . .

If two diaries now exist...

who writes first?

. . .

At first, the answer seems obvious.

Both.

Why shouldn't they?

They're identical.

They trust each other.

The network is healthy.

Nothing suggests trouble.

. . .

One afternoon, two administrators submit score corrections almost simultaneously.

Different candidates.

Different examinations.

Different application servers.

Each server receives one request.

Each opens its own diary.

Each begins writing.

. . .

For a brief moment...

everything appears normal.

Then the network slows.

Not enough to trigger an alert.

Just enough.

The first diary records one correction.

The second diary records the other.

Both continue confidently.

Both believe they are preserving history.

Neither realizes they have just written different first pages of the same chapter.

. . .

Imagine two historians writing the official record of a nation.

One records the coronation before the treaty.

The other records the treaty before the coronation.

Every event happened.

Every sentence is accurate.

Yet generations later...

they describe different histories.

The conflict isn't in the facts.

It's in the order.

***History is defined not only by what happened, but
by what happened first.***

. . .

The engineers gather once more.

Nobody discusses disk performance.

Nobody mentions network bandwidth.

Everyone stares at the same question.

"Who decides the next sentence?"

. . .

Yesterday, multiple diaries protected memory.

Today...

multiple writers threaten it.

The architecture has solved forgetting.

It has accidentally created disagreement.

. . .

Someone proposes merging the diaries afterward.

Reasonable.

Until two updates affect the same candidate.

Which correction wins?

Someone suggests timestamps.

Until two clocks disagree.

Another engineer proposes trusting the fastest machine.

Until the fastest machine isn't the most informed.

Every shortcut quietly depends on another assumption.

Every assumption eventually fails.

. . .

Hours pass.

Whiteboards fill.

Coffee grows cold.

The room becomes strangely calm.

Not because the solution is obvious.

Because the problem has finally become clear.

The architecture does not need faster writers.

It needs one trusted author.

. . .

That sentence changes everything.

Not forever.

Not for every task.

Only for the next piece of shared history.

One voice speaks.

The others listen.

Then they faithfully preserve what was written.

The authority is temporary.

The history is permanent.

. . .

The next deployment introduces no dramatic visual changes.

No impressive dashboard.

No spectacular graph.

The architecture simply behaves differently.

Before history changes...

everyone quietly agrees who may hold the pen.

Only then does the writing begin.

Users never notice.

They shouldn't.

The greatest engineering decisions are often invisible.

. . .

Weeks later, another correction arrives.

Another follows milliseconds later.

This time...

both wait.

One becomes history.

The other follows it.

Not because one request mattered more.

Because tomorrow deserves one answer more than today deserves two.

. . .

As evening falls, you open the notebook you've carried since the first incident.

You flip through its pages.

Connection pools.

Caches.

Retries.

Traffic.

Leadership.

Truth.

Memory.

Every lesson has been preparing for this one realization.

A distributed system can survive slow servers.

Failed disks.

Broken networks.

Even uncertainty.

It cannot survive competing histories.

. . .

Before leaving, you write one final sentence beneath every previous note.

"The hardest part of writing isn't storing information.

It's deciding what the future is allowed to remember."

Tomorrow...

a hundred million readers will arrive.

They will never question the history they're reading.

They will simply trust it.

The architecture has spent this entire journey earning that trust.

ENGINEER'S NOTEBOOK

Today I realized that writing is an act of responsibility.

Every write becomes someone else's reality.

Every reader inherits the decisions of a writer they will never meet.

The purpose of distributed systems isn't to make writing faster.

It's to make history trustworthy.

CHAPTER 25

One Hundred Million Reads

The clock reads 8:59 AM.

The operations room is quiet.

Coffee grows cold beside glowing monitors.

Someone refreshes the dashboard one last time.

Every graph looks ordinary.

Exactly as it should.

• • •

At 9:00 AM...

the first request arrives.

Then another.

Then a thousand.

Then a hundred thousand.

Then millions.

Not gradually.

All at once.

The internet has arrived.

. . .

Months ago, this moment would have terrified you.

One application server.

One database.

One engineer.

You would have watched CPU climb.

Connections disappear.

Queries slow.

Users refresh.

Chaos multiply.

Today...

you watch something else.

Calm.

. . .

The first request reaches the architecture.

It isn't asking for much.

Only one examination result.

One piece of shared history.

The answer begins a journey.

. . .

A familiar copy quietly answers first.

Not because it owns the truth.

Because it remembers it.

The database continues resting.

Exactly as it should.

You smile.

You remember the day you stopped thinking of copies as shortcuts.

And started thinking of them as responsibility.

. . .

The request moves onward.

No application wonders whether another server should answer.

That decision has already been made.

Quietly.

Continuously.

Millions of times.

You remember the chapter where adding another server created another decision.

The decision never disappeared.

It simply became invisible.

. . .

Another request arrives.

Then another.

Ten thousand ask exactly the same question.

Only one performs the work.

The others patiently share the answer.

Months ago...

that idea was called avoiding repeated work.

Today...

it simply feels obvious.

• • •

Far away...

one application server disappears.

Nobody notices.

Another quietly continues the conversation.

Responsibility moves.

History does not.

• • •

A network briefly slows between two machines.

Some requests take slightly longer.

The architecture hesitates before making new decisions.

Not because it is indecisive.

Because certainty deserves patience.

• • •

A downstream service becomes slower.

For a brief moment...

the architecture stops asking it for help.

Not angrily.

Not permanently.

Just long enough for both systems to breathe again.

You remember learning that kindness sometimes looks like silence.

. . .

Elsewhere...

graphs rise.

Logs accumulate.

Metrics change.

Traces grow longer.

No single dashboard explains the morning.

Together...

they tell its story.

You no longer search for *the* truth.

You listen to many witnesses.

. . .

An administrator submits a correction.

The request slows.

Not because the architecture is overloaded.

Because history is changing.

Somewhere...

machines quietly agree.

One trusted voice writes.

Others remember.

Nobody races.

Nobody argues.

Tomorrow has just inherited another fact.

. . .

The correction completes.

Almost instantly...

millions of readers begin seeing it.

Not because millions of writes happened.

Because one trusted write became everyone's shared memory.

The title of the book quietly unfolds before you.

One Write.

One Hundred Million Reads.

. . .

The graphs continue climbing.

Ten million requests.

Twenty million.

Fifty million.

A hundred million.

Nothing dramatic happens.

No heroic incident.

No desperate debugging.

No last-minute optimization.

Only thousands of ordinary decisions...

each learned from an earlier failure.

. . .

You realize something unexpected.

Not one component in today's architecture exists because an engineer admired it.

Every cache remembers a painful lesson.

Every timeout recalls an exhausting incident.

Every retry carries the memory of a previous mistake.

Every refusal protects tomorrow's stability.

Every observer remembers that truth has witnesses.

Every diary remembers that machines forget.

Every leader remembers that history needs one author.

The architecture isn't a collection of technologies.

It's a collection of memories.

. . .

You think back to the first whiteboard.

One application.

One database.

One engineer.

Nothing about that drawing was wrong.

It simply belonged to a younger system.

Growth never demanded brilliance.

Only honesty.

Each new lesson appeared only after the previous architecture could no longer answer reality's questions.

. . .

Outside...

another hundred thousand users arrive.

None of them know what happened over the past months.

They don't know about stale caches.

Or exhausted connections.

Or duplicated reports.

Or uncertain leaders.

Or forgotten history.

They shouldn't.

Great engineering is rarely visible.

Its greatest achievement is making complexity disappear before anyone notices it.

. . .

The dashboard remains calm.

Not because the internet became simpler.

Because the architecture learned.

One inevitable decision at a time.

You close your notebook.

The same notebook you've carried since the first incident.

The final page has only one sentence.

"Readers never see the architecture.

They only inherit its decisions."

The clock reads 9:01 AM.

The internet has arrived.

And the system is ready.

ENGINEER'S NOTEBOOK

Today I finally understood the title.

One write is not a database operation.

It is a promise.

One hundred million reads are not a scaling problem.

They are the reward for keeping that promise.

*Every chapter in this book has been about earning the right to
make that promise once.*

Tomorrow Morning, Revisited

A GUIDED WALK THROUGH ONE REQUEST

Tomorrow has arrived.

Again.

9:00 AM.

The IPO allotment results are now public.

Across the country, millions of people reach for the same answer.

"Did I get the allotment?"

Yesterday, that question looked terrifying.

Today...

it looks familiar.

Let's follow a single request.

One request.

One journey.

. . .

A student opens the website from a hostel room.

A father checks from his office before a meeting.

Someone refreshes the page every few seconds inside a crowded train.

Millions of people.

One question.

"What's my allotment status?"

The internet doesn't see millions of stories.

It only sees millions of requests.

. . .

The request reaches the system.

Months ago, there was only one application server waiting for it.

That server no longer exists.

Not because it failed.

Because success outgrew it.

Now the request arrives at many identical servers.

No server asks,

"Should I answer?"

That decision has already been made.

Quietly.

Long before the request arrived.

You remember why.

One destination made every decision trivial.

Many destinations required judgment.

. . .

The chosen server receives the request.

It could immediately ask the database.

Instead...

it pauses.

Somebody may have already asked this exact question a few milliseconds earlier.

If yesterday's answer is still trustworthy...

why make tomorrow's database remember it again?

The server looks somewhere else first.

Somewhere built to remember repeated questions.

You already know its name.

You also know why it exists.

. . .

The answer is waiting.

No database query.

No expensive computation.

No unnecessary work.

The request continues almost immediately.

Not because the database became faster.

Because the architecture learned that repeating work is rarely a sign of intelligence.

. . .

Another request arrives.

Same PAN number.

Same answer.

Another.

Then another.

Thousands of people unknowingly ask exactly the same question.

The architecture quietly answers them all with the same
remembered truth.

The database never notices.

Exactly as it should.

. . .

Some questions are different.

The answer isn't remembered yet.

This time...

the database speaks.

Not because memory failed.

Because memory knew when to ask for help.

The answer returns.

The architecture remembers it for the next visitor.

Yesterday's effort becomes tomorrow's shortcut.

. . .

Outside...

the requests continue multiplying.

One server begins working harder than the others.

Nobody rushes to fix it.

Nobody even notices.

The next requests simply begin arriving somewhere else.

Responsibility moves.

The question does not.

The user never knows another server answered instead.

They shouldn't.

. . .

One machine suddenly disappears.

Power failure.

Hardware fault.

It doesn't matter.

The architecture quietly stops sending requests there.

No announcement.

No celebration.

Just another decision.

Somewhere else, another server picks up the conversation exactly where the previous one stopped.

Because conversations belong to users.

Not to machines.

. . .

The morning grows busier.

Graphs begin climbing.

Requests per second.

Latency.

Error rates.

Nothing on the dashboard explains the morning.

Each graph remembers one small part of it.

Elsewhere...

logs quietly collect every decision.

Not opinions.

Events.

Traces begin following individual requests through the architecture.

One request.

One journey.

Many witnesses.

You no longer expect one place to explain everything.

Truth has always needed multiple perspectives.

. . .

Suddenly...

one downstream service begins slowing down.

The application notices.

It waits.

Not forever.

Only long enough to remain hopeful.

When hope runs out...

it lets go.

Not because it is impatient.

Because healthy systems know when waiting stops being helpful.

A few moments later...

the architecture tries again.

Carefully.

Optimistically.

The future often behaves differently from the present.

Sometimes it doesn't.

When repeated optimism begins harming everyone else...

the architecture stops asking for a while.

Not as punishment.

As mercy.

Some recoveries need silence before they need traffic.

. . .

Meanwhile...

inside another part of the system...

something far less common happens.

An exchange reports that one allotment calculation was incorrect.

A correction must be made.

Everything changes.

Until now...

every request has been reading history.

This request will change it.

The architecture becomes noticeably quieter.

No rush.

No shortcuts.

History deserves neither.

One trusted writer receives the correction.

Not because it is the smartest machine.

Not because it is the fastest.

Because tomorrow deserves one story more than it deserves two.

The correction is written.

Others remember it.

Only then...

does the architecture allow the new truth to spread.

A million future readers will never know this moment happened.

They will simply receive the correct answer.

That quiet confidence is the entire purpose of the architecture.

. . .

The morning continues.

Ten million requests.

Twenty million.

Fifty million.

One hundred million.

The dashboards remain ordinary.

Support engineers remain seated.

Users continue checking their allotment status.

Nobody congratulates the cache.

Nobody thanks the servers.

Nobody notices the careful decisions happening every millisecond beneath a simple web page.

That is success.

Engineering succeeds most completely when nobody realizes it was difficult.

. . .

You close your notebook.

The same notebook you opened before the first request of this story.

It no longer contains technologies.

It contains decisions.

Every page answers one question.

Why did this become necessary?

You finally understand something that seemed impossible on the first page of this book.

The architecture serving today's IPO allotment system was never designed in one afternoon.

It was written over years.

One incident.

One lesson.

One careful decision at a time.

And tomorrow...

when another engineer inherits this system...

they won't inherit Redis.

Or load balancers.

Or retries.

Or leaders.

They'll inherit something far more valuable.

The reasons those things exist.

That is the only architecture worth preserving.

Epilogue

TOMORROW ARRIVED

The clock reads 9:02 AM.

The morning everyone feared has already begun.

The dashboards remain calm.

Support tickets never arrive.

Nobody rushes toward the operations room.

Nobody cancels meetings.

Nobody notices.

Which is exactly how success is supposed to feel.

. . .

Outside the data center...

one hundred million people continue reading.

Some search for examination results.

Some refresh a page.

Some retry after losing internet connectivity.

Some never realize how many decisions quietly happened before a single answer reached their screen.

They simply receive a response.

Correct.

Consistent.

Ordinary.

. . .

Inside the architecture...

nothing extraordinary happens either.

A cache quietly remembers.

A server quietly accepts responsibility.

Another quietly steps aside.

A request quietly waits.

Another quietly gives up.

A diary quietly grows by one page.

A new server quietly learns yesterday's story.

History quietly survives another morning.

No component demands attention.

Each simply performs the responsibility it earned.

. . .

Months ago...

none of this existed.

There was no cache.

No traffic distribution.

No shared memory.

No observers.

No leaders.

No diary.

There didn't need to be.

The simplest architecture was the correct architecture.

The internet simply wasn't asking harder questions yet.

. . .

People often imagine distributed systems as something designed all at once.

A grand blueprint.

A complicated diagram.

An architect predicting every future challenge before the first line of code is written.

Reality is quieter.

Much quieter.

One incident.

One lesson.

One careful decision.

Repeated over years.

. . .

Every chapter in this story began with the same sentence.

"Something unexpected happened."

Not because engineers lacked knowledge.

Because reality always knows something we don't.

Engineering has never been the art of predicting every problem.

It has always been the discipline of responding honestly when prediction fails.

. . .

Along the way...

we stopped thinking about technology.

Caches became memory.

Retries became hope.

Traffic distribution became judgment.

Leadership became responsibility.

Logs became diaries.

Consensus became humility.

Writing became a promise.

The names mattered less than the ideas beneath them.

Technologies change.

Principles survive.

. . .

One day...

someone else will inherit this system.

They'll see Redis.

PostgreSQL.

Traffic distribution.

Replication.

Observers.

Leaders.

They may wonder why each piece exists.

I hope they never remove one simply because it looks unnecessary.

Every component is a scar.

Every scar remembers a lesson.

. . .

Years from now...

new technologies will replace today's.

New databases.

New programming languages.

New deployment platforms.

Perhaps even entirely new ways of building software.

It won't matter.

Reality will continue asking the same questions.

Can memory be trusted?

Can uncertainty be tolerated?

Can history survive failure?

Can millions of strangers believe the same answer?

Every generation of engineers will invent different tools.

They will rediscover the same principles.

. . .

You close the notebook.

The notebook is full now.

Connection pools.

Copies.

Retries.

Observers.

Leaders.

Memory.

History.

Every page describes a problem that once seemed impossible.

Every page also describes a system that quietly became wiser
because of it.

You place the notebook back on the shelf.

Tomorrow...

someone else will open it.

Not because they want to build a distributed system.

Because reality has begun asking them the same questions.

. . .

The architecture outside continues serving requests.

Quietly.

Patiently.

Confidently.

The internet has no idea how many lessons are hidden beneath a
single response.

It never should.

That is the privilege.

And the responsibility.

Of engineering.

. . .

This book began with one question.

"Will the system survive tomorrow?"

It ends with a different one.

"What will tomorrow teach the system next?"

Because no distributed system is ever finished.

No architecture is ever complete.

No final diagram remains final for long.

There will always be another incident.

Another lesson.

Another careful decision.

Another page in the diary.

And that is exactly how it should be.

. . .

If there is one idea I hope remains long after the names of today's technologies have faded, it is this:

Distributed systems are not monuments to complexity.


They are the accumulated memory of every problem we refused to forget.


The Journey at a Glance


Every architecture begins with confidence.

Every architecture ends with humility.

What happened between those two moments is the story you just read.







At no point did the architecture become complicated because engineers wanted it to.

Every new idea appeared because the previous one could no longer answer reality.

That is how distributed systems actually evolve.

One Sentence from Every Chapter

Prologue

Tomorrow always arrives.

• • •

Chapter 1 — The Simplest Backend You'll Ever Build

Every architecture begins with the simplest answer it can honestly justify.

• • •

Chapter 2 — Build the Wrong System First

The first version is not supposed to last; it is supposed to teach.

• • •

Chapter 3 — Nothing Is Broken. So Why Is Everyone Complaining?

Users experience journeys, not services.

• • •

Chapter 4 — PostgreSQL Didn't Fail. You Did.

Every technology enters your architecture carrying an invoice.

• • •

Chapter 5 — The Most Expensive Question in Engineering

The fastest work is the work you never repeat.

• • •

Chapter 6 — The Fastest Database Query Costs Zero

A cache moves complexity from computation to correctness.

• • •

Chapter 7 — Every Cache Eventually Lies

Every copy is a promise to stay consistent with the truth.

• • •

Chapter 8 — Your Biggest DDoS Attack Doesn't Know It's Attacking You

Humans create load more creatively than machines ever will.

• • •

Chapter 9 — When One Server Becomes Two

Scale is not multiplication; it is coordination.

• • •

Chapter 10 — The Decision Maker Nobody Notices

Every routing decision is an opinion about the future.

• • •

Chapter 11 — Yesterday's Truth

The present is distributed.

• • •

Chapter 12 — The System Starts Watching Itself

Automation is architecture acting on its own beliefs.

• • •

Chapter 13 — Four Witnesses, One Incident

Every observability signal is a witness, not the truth.

• • •

Chapter 14 — The Waiting Game

Waiting is work that nobody notices until it disappears.

• • •

Chapter 15 — Hope Is Not a Strategy

Every retry assumes the future will be different.

• • •

Chapter 16 — Learning to Say No

Sometimes the healthiest response is refusal.

• • •

Chapter 17 — When Everyone Decides

Leadership exists to make responsibility unambiguous.

• • •

Chapter 18 — The Day Truth Split in Two

Truth becomes fragile the moment distance enters the conversation.

• • •

Chapter 19 — When Nobody Gets to Decide Alone

Sometimes refusing to decide is the only honest decision.

• • •

Chapter 20 — The Memory Everyone Must Share

Reads consume history. Writes create it.

• • •

Chapter 21 — The Half-Written Story

A write that cannot survive failure has not truly happened.

• • •

Chapter 22 — The Diary

Memory becomes trustworthy only after it leaves memory.

• • •

Chapter 23 — The Cost of Forgetting

*History becomes durable only after it can survive
being forgotten by one place.*

• • •

Chapter 24 — One Write

*The hardest part of writing is deciding what the future
is allowed to remember.*

• • •

Chapter 25 — One Hundred Million Reads

*Readers never see the architecture. They only inherit
its decisions.*

• • •

Epilogue — Tomorrow Arrived

*Distributed systems are not monuments to complexity.
They are the accumulated memory of every problem
we refused to forget.*

ABOUT THE AUTHOR

Kush Mishra is a backend software engineer specializing in distributed systems, event-driven architectures and large-scale backend infrastructure. Throughout his career, he has built and operated systems that process tens of millions of events per hour, led zero-downtime migrations of mission-critical services, and designed software that continues to operate reliably under massive scale.

Outside his day-to-day engineering work, Kush enjoys exploring the deeper principles behind how distributed systems evolve from the first monolith to globally distributed architectures. He believes that the most valuable engineering lessons are rarely about individual technologies, but about understanding the trade-offs that shape systems over time.

He has been recognized through the Meta Bug Bounty Program for identifying a security vulnerability via reverse engineering and is an ACM ICPC Regionalist.

Kush graduated from Jadavpur University.

One Write. One Hundred Million Reads. is his first book.